

4__ML__Global

September 3, 2025

1 Modeling Daily Passenger Counts with Machine Learning

The goal of this notebook is to model and compare different Machine Learning approaches using the `global_daily_df` dataset.

This dataset aggregates daily passenger counts (boardings and alightings) across the entire TPG network.

The objectives are:

- to evaluate the ability of different models to capture temporal and contextual patterns in the data,
- to compare their performance using standard metrics (MAE, RMSE, R^2),

This work will provide insights into the strengths and limitations of several Machine Learning models when applied to public transport demand forecasting.

1.0.1 Method :

We will implement and compare several Machine Learning models, using a consistent pipeline for preprocessing and evaluation.

1.0.2 Remark :

The analysis is limited to daily aggregated data.
Stop-level details are addressed in the next notebook.

```
[1]: # libraries
      %matplotlib inline
      import numpy as np
      import pandas as pd
      import matplotlib.pyplot as plt
      import seaborn as sns
```

2 1) Loading the data

```
[2]: global_daily_df = pd.read_csv("global_daily_df.csv", sep=",")
      global_daily_df['date'] = pd.to_datetime(global_daily_df['date'])
      display(global_daily_df.info())
```

```

<class 'pandas.core.frame.DataFrame'>
RangeIndex: 1415 entries, 0 to 1414
Data columns (total 18 columns):
#   Column                                Non-Null Count  Dtype
---  -
0   date                                  1415 non-null   datetime64[ns]
1   nb_de_montees                        1415 non-null   float64
2   nb_de_descentes                      1415 non-null   float64
3   frequentation_totale                 1415 non-null   float64
4   frequentation_totale_lag_1           1415 non-null   float64
5   frequentation_totale_lag_7           1415 non-null   float64
6   rolling_mean_7d                     1415 non-null   float64
7   weather_code                        1415 non-null   int64
8   log_frequentation_totale             1415 non-null   float64
9   log_frequentation_totale_lag_1       1415 non-null   float64
10  log_frequentation_totale_lag_7       1415 non-null   float64
11  log_rolling_mean_7d                  1415 non-null   float64
12  indice_jour_semaine                 1415 non-null   int64
13  indice_semaine                      1415 non-null   int64
14  mois                                1415 non-null   int64
15  annee                                1415 non-null   int64
16  est_weekend                         1415 non-null   int64
17  est_ferie                           1415 non-null   int64
dtypes: datetime64[ns](1), float64(10), int64(7)
memory usage: 199.1 KB

None

```

```

[3]: print(f"The data covers the date range from {global_daily_df['date'].min()} to_
      ↳{global_daily_df['date'].max()}")

```

The data covers the date range from 2021-08-08 00:00:00 to 2025-06-22 00:00:00

3 2) Splitting the data

3.0.1 Observation :

For a time series, the chronological order must be preserved — we do not shuffle the rows.
The dataset should be split sequentially into **train** → **validation** → **test**.

3.0.2 Method :

Two approaches can be used:

1. **Fixed split** by defining explicit temporal boundaries.
2. **TimeSeriesSplit for cross-validation**, which allows splitting the training set into several train/validation folds while preserving temporal order.

3.1 a) Fixed split

```
[4]: # Définition des bornes temporelles
train_end = pd.to_datetime('2023-12-31')
valid_start = pd.to_datetime('2024-01-01')
valid_end = pd.to_datetime('2024-06-30')
test_start = pd.to_datetime('2024-07-01')
test_end = pd.to_datetime('2025-06-22')

def split_dataset(dataset):
    train_df = dataset[dataset['date'] <= train_end]
    valid_df = dataset[(dataset['date'] >= valid_start) & (dataset['date'] <=
    ↪valid_end)]
    test_df = dataset[(dataset['date'] >= test_start) & (dataset['date'] <=
    ↪test_end)]
    return train_df, valid_df, test_df

# Découpage de global_daily_df
global_train_df, global_valid_df, global_test_df =
    ↪split_dataset(global_daily_df)

# Découpage de TPG_meteo_all_df
# tpg_train_df, tpg_valid_df, tpg_test_df = split_dataset(TPG_meteo_all_df)

# Affichage des tailles pour vérification
print("GLOBAL DAILY SPLIT:")
print(f"Train: {len(global_train_df)} jours")
print(f"Valid: {len(global_valid_df)} jours")
print(f"Test : {len(global_test_df)} jours\n")

# print("TPG + METEO SPLIT:")
# print(f"Train: {len(tpg_train_df)} lignes")
# print(f"Valid: {len(tpg_valid_df)} lignes")
# print(f"Test : {len(tpg_test_df)} lignes")
```

GLOBAL DAILY SPLIT:

Train: 876 jours


Valid: 182 jours

Test : 357 jours

```
[5]: # Plot training and test sets
fig = plt.figure(figsize=(14, 4))
plt.plot(global_train_df.frequentation_totale, label="Train set")
plt.plot(global_valid_df.frequentation_totale, label="Validation set",
    ↪color="g")
```

```
plt.plot(global_test_df.frequentation_totale, label="Test set", color="r")

plt.title("train, validation and test sets")
plt.ylabel("loads (MW)")
plt.legend()
plt.show()
```



3.2 b) TimeSeriesSplit for cross-validation

```
[6]: # load the library
from sklearn.model_selection import TimeSeriesSplit

train = global_daily_df.loc[global_daily_df['date'] <= "2024-06-30"].
    ↪reset_index(drop=True)
test = global_daily_df.loc[global_daily_df['date'] >= "2024-07-01"].
    ↪reset_index(drop=True)

# Split train set into train/validation sets
cv = TimeSeriesSplit(n_splits=5)

[7]: # Plot iterations of the training and validation sets

# Création du graphique
fig, axes = plt.subplots(5, 1, figsize=(14, 10), sharex=True, sharey=True)
plt.suptitle("TimeSeriesSplit - Visualisation des splits Train/Validation",
    ↪fontsize=16)

for i, (tr_idx, val_idx) in enumerate(cv.split(train)):
    ax = axes[i]

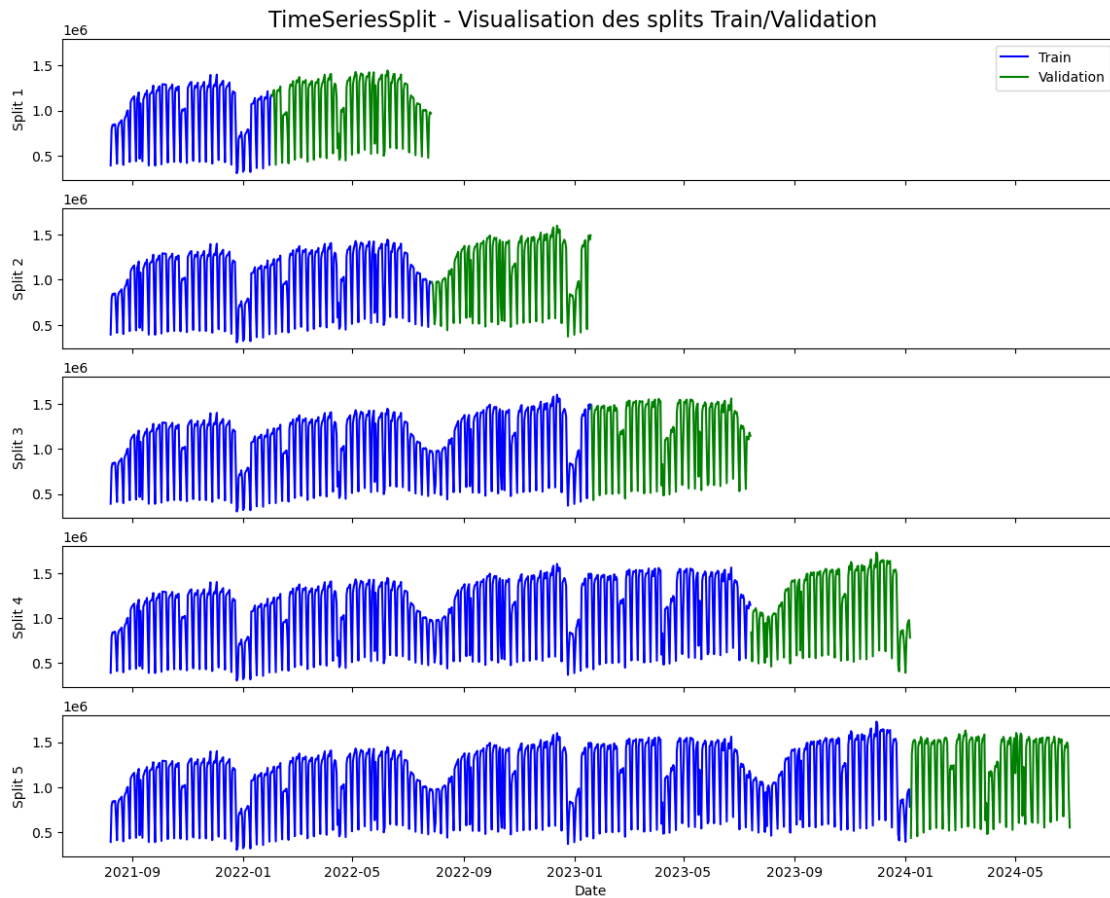
    # Affichage avec dates en abscisse
```

```

    ax.plot(train["date"].iloc[tr_idx], train["frequentation_totale"].
    ↪iloc[tr_idx], label="Train", color="blue")
    ax.plot(train["date"].iloc[val_idx], train["frequentation_totale"].
    ↪iloc[val_idx], label="Validation", color="green")
    ax.set_ylabel(f"Split {i+1}")

axes[-1].set_xlabel("Date")
axes[0].legend()
plt.subplots_adjust(top=0.95)
plt.show()

```



3.2.1 Conclusion :

Since the data are continuous over the entire period, the `global_daily_df` dataset will be split using **TimeSeriesSplit**.

This approach allows the model to be trained and validated across 5 sequential cycles, while preserving temporal order.

```
[8]: # Liste des features (version log pour les séries temporelles)
```

```
features = [
    'log_frequentation_totale_lag_1',
    'log_frequentation_totale_lag_7',
    'log_rolling_mean_7d',
    'indice_jour_semaine',
    'indice_semaine',
    'mois',
    'annee',
    'est_weekend',
    'est_ferie'
]
```

```
target = 'log_frequentation_totale'
```

```
# Définition des jeux de données
```

```
X_tr = train[features]
```

```
y_tr = train[target]
```

```
X_te = test[features]
```

```
y_te = test[target]
```

```
X_tr.head()
```

```
[8]:
```

	log_frequentation_totale_lag_1	log_frequentation_totale_lag_7	\
0	13.202760	12.834165	
1	12.878287	13.522369	
2	13.581051	13.589748	
3	13.642689	13.582402	
4	13.634925	13.612583	

	log_rolling_mean_7d	indice_jour_semaine	indice_semaine	mois	annee	\
0	13.461451	6	31	8	2021	
1	13.470585	0	32	8	2021	
2	13.479297	1	32	8	2021	
3	13.487801	2	32	8	2021	
4	13.494145	3	32	8	2021	

	est_weekend	est_ferie
0	1	0
1	0	0
2	0	0
3	0	0
4	0	0

4 3) Baseline model

We begin with a very simple model: using only the mean value.

```
[9]: from sklearn.metrics import mean_squared_error, mean_absolute_error, r2_score

def compute_statistics(model_name, y_te, y_te_pred_log) :
    mae_log = mean_absolute_error(y_te, y_te_pred_log)
    rmse_log = np.sqrt(mean_squared_error(y_te, y_te_pred_log))
    r2_log = r2_score(y_te, y_te_pred_log)

    print(f"{model_name} (log space):")
    print(f"  MAE : {mae_log:.4f}")
    print(f"  RMSE: {rmse_log:.4f}")
    print(f"  R2  : {r2_log:.4f}")

    y_te_pred_real = np.exp(y_te_pred_log)
    y_te_real = np.exp(y_te)

    mae_real = mean_absolute_error(y_te_real, y_te_pred_real)
    rmse_real = np.sqrt(mean_squared_error(y_te_real, y_te_pred_real))
    r2_real = r2_score(y_te_real, y_te_pred_real)

    print("\n")
    print(f"{model_name} (real space):")
    print(f"  MAE : {mae_real:.4f}")
    print(f"  RMSE: {rmse_real:.4f}")
    print(f"  R2  : {r2_real:.4f}")

    return mae_log, rmse_log, r2_log, mae_real, rmse_real, r2_real
```

```
[10]: from sklearn.metrics import mean_squared_error, mean_absolute_error, r2_score

y_te_pred_log = np.median(y_tr) * np.ones(len(y_te))
#y_te_pred_real = np.median(y_tr_real) * np.ones(len(y_te_real))

base_model_name = "Baseline (median)"
base_mae_log, base_rmse_log, base_r2_log, base_mae_real, base_rmse_real,   
↪base_r2_real = compute_statistics(base_model_name, y_te, y_te_pred_log)
```

Baseline (median) (log space):

MAE : 0.3015
RMSE: 0.3827
R² : -0.0248

Baseline (median) (real space):

MAE : 331204.8508

RMSE: 378400.1937

R^2 : -0.0001

5 4) Ridge regression: linear model with regularization

Ridge regression is a regularized linear model that penalizes large coefficients.

This reduces the impact of extreme values or highly correlated variables, and helps to limit overfitting.

```
[11]: from sklearn.pipeline import Pipeline
      from sklearn.preprocessing import OneHotEncoder
      from sklearn.compose import ColumnTransformer
      from sklearn.preprocessing import StandardScaler

      cat_features = ['mois', 'indice_jour_semaine', 'annee']
      num_features = [
          'log_frequentation_totale_lag_1',
          'log_frequentation_totale_lag_7',
          'log_rolling_mean_7d',
          'est_weekend',
          'est_ferie'
      ]

      # Define the transformer
      cat_transformer = Pipeline([("onehot", OneHotEncoder(handle_unknown='ignore'))])
      num_transformer = Pipeline([("scaler", StandardScaler())])

      # Apply
      preprocessor = ColumnTransformer(
          transformers=[
              ('cat', cat_transformer, cat_features),
              ('num', num_transformer, num_features)
          ]
      )
```

```
[12]: from sklearn.linear_model import Ridge
      from sklearn.model_selection import GridSearchCV

      # Create the pipeline
      pipe_ridge = Pipeline([("preprocessor", preprocessor), ("ridge", Ridge())])

      # Create cross-validation object
      cv = TimeSeriesSplit(n_splits=5) # 1 month validation , test_size=182

      # Create grid for alpha
```



```

grid = {"ridge__alpha": np.logspace(-4, 4, num=20)}

# Create the grid search object
model_ridge = GridSearchCV(
    pipe_ridge,
    grid,
    cv=cv,
    return_train_score=True,
    scoring="neg_mean_absolute_error",
    verbose=1,
)

```

```

[13]: # Fit on the training set
model_ridge.fit(X_tr, y_tr)

# Evaluate on the test set
y_te_pred_log = model_ridge.predict(X_te)

ridge_model_name = "ridge regression"
ridge_mae_log, ridge_rmse_log, ridge_r2_log, ridge_mae_real, ridge_rmse_real, u
↪ridge_r2_real = compute_statistics(ridge_model_name, y_te, y_te_pred_log)

```

Fitting 5 folds for each of 20 candidates, totalling 100 fits

ridge regression (log space):

MAE : 0.0702

RMSE: 0.1362

R^2 : 0.8702

ridge regression (real space):

MAE : 75017.5051

RMSE: 130343.4130

R^2 : 0.8813

```

[14]: # Collect results in a DataFrame
cv_results = pd.DataFrame(model_ridge.cv_results_)

# Print a few interesting columns
cols = [
    "mean_test_score",
    "std_test_score",
    "mean_train_score",
    "std_train_score",
    "param_ridge__alpha",
]
cv_results[cols].sort_values("mean_test_score", ascending=False).head(5)

```

```
[14]:
```


	mean_test_score	std_test_score	mean_train_score	std_train_score	\
12	-0.079444	0.006534	-0.071999	0.000692	
11	-0.080861	0.011062	-0.070305	0.001233	
13	-0.081899	0.006565	-0.078118	0.003464	
10	-0.082904	0.015712	-0.069892	0.001791	
9	-0.083877	0.017800	-0.069796	0.002036	

	param_ridge__alpha
12	11.288379
11	4.281332
13	29.763514
10	1.623777
9	0.615848

```
[15]: # Plot train and validation curves
plt.semilogx(
    cv_results["param_ridge__alpha"],
    -cv_results["mean_train_score"],
    label="training error",
)
plt.semilogx(
    cv_results["param_ridge__alpha"],
    -cv_results["mean_test_score"],
    label="validation error",
)

# Add marker for best score
plt.scatter(
    model_ridge.best_params_.values(),
    -1 * model_ridge.best_score_,
    marker="x",
    c="red",
    zorder=10,
)
plt.xlabel("$alpha$")
plt.ylabel("MAE")
plt.title("Evolution of training/validation errors with $alpha$")

plt.legend()
plt.show()
```



5.0.1 Observation :


The MAE remains stable for small values of alpha and starts to increase significantly when alpha becomes too large.

The optimal performance is reached around alpha = 10, where the validation error is minimized.

```
[16]: # Revenir à l'échelle réelle
y_valid_real = np.exp(y_te)
y_pred_real = np.exp(y_te_pred_log)

plt.figure(figsize=(10, 6))
plt.plot(y_valid_real.values, label='Actual value', linewidth=2)
plt.plot(y_pred_real, label='Prediction', linestyle='--')
plt.title('Comparison between actual values and predictions (validation set)',
          fontsize=14)
plt.xlabel('Days')
plt.ylabel('Ridership')
plt.legend()
```

```
plt.grid(True)
plt.tight_layout()
plt.show()
```



5.0.2 Observation :

The model captures the general dynamics of daily ridership well, including seasonal variations and weekly patterns.

However, discrepancies remain during some peaks and troughs, where the predictions tend to underestimate or overestimate the actual values.

6 5) Analyzing the influence of weather

Does weather have an influence on passenger counts?

We have seen that the correlation between weather and ridership is very weak.

The question is whether this characteristic remains true when using a model such as Ridge regression.

```
[17]: # Liste des features (version log pour les séries temporelles)
features = [
    'log_frequentation_totale_lag_1',
    'log_frequentation_totale_lag_7',
    'log_rolling_mean_7d',
    'indice_jour_semaine',
    'indice_semaine',
```

```

    'mois',
    'annee',
    'est_weekend',
    'est_ferie',
    'weather_code'
]

target = 'log_frequentation_totale'

# Définition des jeux de données
X_tr = train[features]
y_tr = train[target]

X_te = test[features]
y_te = test[target]

X_tr.head()

cat_features = ['mois', 'indice_jour_semaine', 'annee', 'weather_code']
num_features = [
    'log_frequentation_totale_lag_1',
    'log_frequentation_totale_lag_7',
    'log_rolling_mean_7d',
    'est_weekend',
    'est_ferie'
]

# Define the transformer
cat_transformer = Pipeline([("onehot", OneHotEncoder(handle_unknown='ignore'))])
num_transformer = Pipeline([("scaler", StandardScaler())])

# Apply
preprocessor = ColumnTransformer(
    transformers=[
        ('cat', cat_transformer, cat_features),
        ('num', num_transformer, num_features)
    ]
)

# Create the pipeline
pipe_ridge = Pipeline([("preprocessor", preprocessor), ("ridge", Ridge())])

# Create cross-validation object
cv = TimeSeriesSplit(n_splits=5) # 1 month validation , test_size=182

# Create grid for alpha

```

```

grid = {"ridge__alpha": np.logspace(-4, 4, num=20)}

# Create the grid search object
model_ridge = GridSearchCV(
    pipe_ridge,
    grid,
    cv=cv,
    return_train_score=True,
    scoring="neg_mean_absolute_error",
    verbose=1,
)

# Fit on the training set
model_ridge.fit(X_tr, y_tr)

# Evaluate on the test set
y_te_pred_log = model_ridge.predict(X_te)

ridge_w_model_name = "ridge with weather"
ridge_w_mae_log, ridge_w_rmse_log, ridge_w_r2_log, ridge_w_mae_real,
↪ridge_w_rmse_real, ridge_w_r2_real = compute_statistics(ridge_w_model_name,
↪y_te, y_te_pred_log)

```

Fitting 5 folds for each of 20 candidates, totalling 100 fits

ridge with weather (log space):

MAE : 0.0706
 RMSE: 0.1362
 R^2 : 0.8701

ridge with weather (real space):

MAE : 75725.2624
 RMSE: 130544.9033
 R^2 : 0.8810

Variante	RMSE (log)	MAE (log)	R^2	RMSE (réel)	MAE (réel)	R^2
Baseline	0.3827	0.3015	-0.0248	378400.1933	331204.8454	-0.0001
ridge sans météo	0.1362	0.0702	0.8702	130343.4130	75017.5051	0.8813
ridge with weather_code	0.1362	0.0706	0.8701	130544.9033	75725.2624	0.8810

6.0.1 Observation :

Including weather features does not produce any significant improvement in the model's performance.

6.0.2 Question :

Which features are the most influential?

6.0.3 Action :

Analyze feature importance to identify the variables that contribute the most to the prediction.

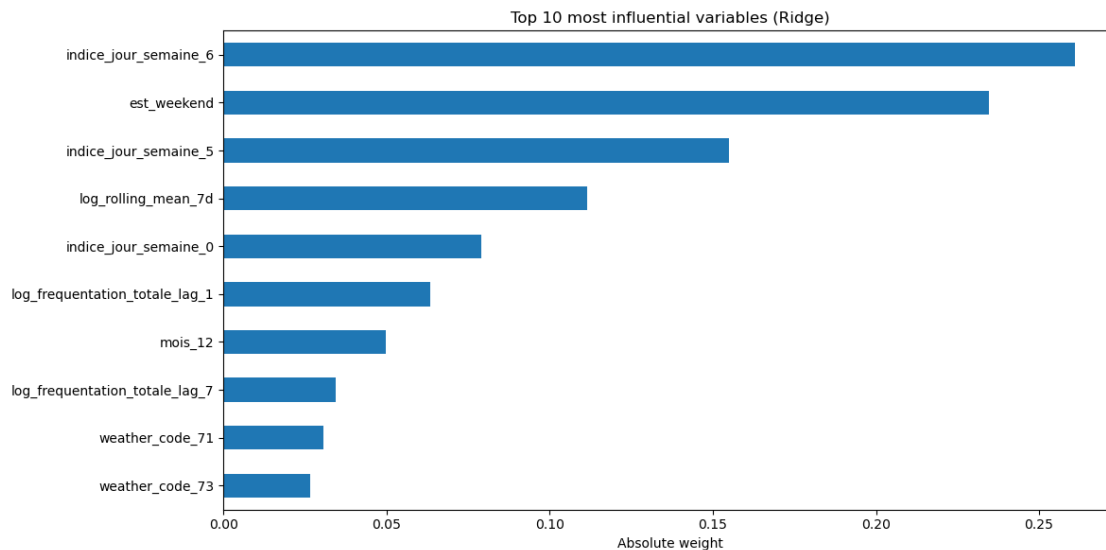
```
[18]: # code by chatGPT
# Récupération des coefficients
coefs = model_ridge.best_estimator_.named_steps['ridge'].coef_

# Récupération des noms de features après preprocessing
onehot_feature_names = model_ridge.best_estimator_.named_steps['preprocessor']\
    .transformers_[0][1].named_steps['onehot']\
    .get_feature_names_out(cat_features)

all_feature_names = np.concatenate([onehot_feature_names, num_features])

# Création d'une série triée par valeur absolue
coef_series = pd.Series(coefs, index=all_feature_names)
coef_series_abs = coef_series.abs().sort_values(ascending=False)

coef_series_abs.head(10).sort_values().plot(
    kind='barh', figsize=(12, 6), title="Top 10 most influential variables_
↳(Ridge)"
)
plt.xlabel("Absolute weight")
plt.tight_layout()
plt.show()
```



6.0.4 Observation :

- The most influential variables are linked to the **day of the week** (`indice_jour_semaine_6`, `indice_jour_semaine_5`, `indice_jour_semaine_0`) and the **weekend indicator** (`est_weekend`).
- Temporal features such as the **7-day rolling mean** and lagged passenger counts (`log_frequentation_totale_lag_1`, `log_frequentation_totale_lag_7`) also play an important role.
- Weather variables (`weather_code_71`, `weather_code_73`) appear in the Top 10 but with relatively low influence.

6.0.5 Conclusion :

Weekly seasonality and past ridership values are the dominant drivers in the Ridge model, while the impact of weather remains marginal.

7 6) Decision Trees

Decision Trees are non-linear models that split the feature space into regions using recursive rules. They can capture interactions and non-linear effects without feature scaling.

```
[19]: from sklearn.tree import DecisionTreeRegressor

# Liste des features (version log pour les séries temporelles)
features = [
    'log_frequentation_totale_lag_1',
    'log_frequentation_totale_lag_7',
    'log_rolling_mean_7d',
    'indice_jour_semaine',
    'indice_semaine',
    'mois',
    'annee',
    'est_weekend',
    'est_ferie',
    'weather_code'
]

target = 'log_frequentation_totale'

# Définition des jeux de données
X_tr = train[features]
y_tr = train[target]

X_te = test[features]
y_te = test[target]
```



```

X_tr.head()

cat_features = ['mois', 'indice_jour_semaine', 'annee', 'weather_code']
num_features = [
    'log_frequentation_totale_lag_1',
    'log_frequentation_totale_lag_7',
    'log_rolling_mean_7d',
    'est_weekend',
    'est_ferie'
]

# Define the transformer
cat_transformer = Pipeline([("onehot", OneHotEncoder(handle_unknown='ignore'))])
num_transformer = Pipeline([("scaler", StandardScaler())])

# Apply
preprocessor = ColumnTransformer(
    transformers=[
        ('cat', cat_transformer, cat_features),
        ('num', num_transformer, num_features)
    ]
)

# Create the pipeline
pipe_ridge = Pipeline([("preprocessor", preprocessor), ("decisionTree",
    ↪ DecisionTreeRegressor(random_state=42))])

# Create cross-validation object
cv = TimeSeriesSplit(n_splits=5) # 1 month validation , test_size=182

# Create grid for alpha
grid = {"decisionTree__max_depth": [3, 5, 7, 8, 9, 10, 11, 12, 15, 20, None] }

# Create the grid search object
model = GridSearchCV(
    pipe_ridge,
    grid,
    cv=cv,
    return_train_score=True,
    scoring="neg_mean_absolute_error",
    verbose=1,
)

# Fit on the training set
model.fit(X_tr, y_tr)

# Evaluate on the test set

```

```
y_te_pred_log = model.predict(X_te)
```

Fitting 5 folds for each of 11 candidates, totalling 55 fits

```
[20]: #display(model.cv_results_)

# Collect results in a DataFrame
cv_results = pd.DataFrame(model.cv_results_)

# Print a few interesting columns
cols = [
    "mean_test_score",
    "std_test_score",
    "mean_train_score",
    "std_train_score",
    "param_decisionTree__max_depth",
]
cv_results[cols].sort_values("mean_test_score", ascending=False).head(5)
```

```
[20]:
```

	mean_test_score	std_test_score	mean_train_score	std_train_score	\
5	-0.096329	0.030621	-0.009042	0.004387	
4	-0.096788	0.029723	-0.014471	0.006790	
9	-0.097480	0.029137	-0.000020	0.000024	
8	-0.098462	0.026859	-0.000951	0.000602	
10	-0.098922	0.028093	0.000000	0.000000	

	param_decisionTree__max_depth
5	10
4	9
9	20
8	15
10	None

```
[21]: # Plot train and validation curves
plt.semilogx(
    cv_results["param_decisionTree__max_depth"],
    -cv_results["mean_train_score"],
    label="training error",
)
plt.semilogx(
    cv_results["param_decisionTree__max_depth"],
    -cv_results["mean_test_score"],
    label="validation error",
)


# Add marker for best score
plt.scatter(
```

```

model.best_params_.values(),
-1 * model.best_score_,
marker="x",
c="red",
zorder=10,
)
plt.xlabel("max_depth")
plt.ylabel("MAE")
plt.title("Evolution of training/validation errors with max_depth")

plt.legend()
plt.show()

```



Observation

We observe a continuous decrease in training error as `max_depth` increases, indicating that the tree is learning more and more details.

The validation error reaches a minimum around `max_depth` 10, then stabilizes, suggesting that deeper trees do not bring further improvement and may lead to overfitting.

Conclusion

The optimal choice of `max_depth` is around 10, providing a good balance between bias and variance.

```
[22]: # model.best_params_

# Fit on the training set
# model.fit(X_tr, y_tr)
# GridSearchCV réentraîne automatiquement un modèle final avec ces best_params_
  ↳ sur tout le jeu X_tr

# Evaluate on the test set
y_te_pred_log = model.predict(X_te)

dt_model_name = "decision tree"
dt_mae_log, dt_rmse_log, dt_r2_log, dt_mae_real, dt_rmse_real, dt_r2_real =
  ↳ compute_statistics(dt_model_name, y_te, y_te_pred_log)
```

decision tree (log space):

```
MAE : 0.0653
RMSE: 0.1456
R2  : 0.8516
```

decision tree (real space):

```
MAE : 68430.1350
RMSE: 146487.1440
R2  : 0.8501
```

8 7) Random Forest

Random Forest is an ensemble method that averages multiple decision trees to reduce variance and improve generalization.

It should mitigate the overfitting we observed with single Decision Trees.

```
[23]: from sklearn.ensemble import RandomForestRegressor

# Liste des features (version log pour les séries temporelles)
features = [
    'log_frequentation_totale_lag_1',
    'log_frequentation_totale_lag_7',
    'log_rolling_mean_7d',
    'indice_jour_semaine',
    'indice_semaine',
    'mois',
    'annee',
    'est_weekend',
    'est_ferie',
    'weather_code'
```

```

]

target = 'log_frequentation_totale'

# Définition des jeux de données
X_tr = train[features]
y_tr = train[target]

X_te = test[features]
y_te = test[target]

X_tr.head()

cat_features = ['mois', 'indice_jour_semaine', 'annee', 'weather_code']
num_features = [
    'log_frequentation_totale_lag_1',
    'log_frequentation_totale_lag_7',
    'log_rolling_mean_7d',
    'est_weekend',
    'est_ferie'
]

# Define the transformer
cat_transformer = Pipeline([("onehot", OneHotEncoder(handle_unknown='ignore'))])
num_transformer = Pipeline([("scaler", StandardScaler())])

# Apply
preprocessor = ColumnTransformer(
    transformers=[
        ('cat', cat_transformer, cat_features),
        ('num', num_transformer, num_features)
    ]
)

# Create the pipeline
pipe_ridge = Pipeline([("preprocessor", preprocessor), ("rfr",
    ↪ RandomForestRegressor(random_state=42, n_jobs=-1))])

# Create cross-validation object
cv = TimeSeriesSplit(n_splits=5) # 1 month validation , test_size=182

# Create grid for alpha
grid = {"rfr__n_estimators": [3, 5, 7, 8, 9, 10, 11, 12, 15, 20] }

# Create the grid search object
model = GridSearchCV(
    pipe_ridge,

```

```

    grid,
    cv=cv,
    return_train_score=True,
    scoring="neg_mean_absolute_error",
    verbose=1,
)

# Fit on the training set
model.fit(X_tr, y_tr)

# Evaluate on the test set
y_te_pred_log = model.predict(X_te)

```

Fitting 5 folds for each of 10 candidates, totalling 50 fits

```

[24]: # display(model.cv_results_)

# Collect results in a DataFrame
cv_results = pd.DataFrame(model.cv_results_)

# Print a few interesting columns
cols = [
    "mean_test_score",
    "std_test_score",
    "mean_train_score",
    "std_train_score",
    "param_rfr__n_estimators",
]
cv_results[cols].sort_values("mean_test_score", ascending=False).head(5)

```

```

[24]:  mean_test_score  std_test_score  mean_train_score  std_train_score  \
2         -0.082408         0.015472         -0.026857         0.001876
9         -0.083800         0.020017         -0.024483         0.001186
1         -0.083848         0.015843         -0.028327         0.002063
8         -0.083919         0.017706         -0.025504         0.002033
3         -0.084982         0.018365         -0.026572         0.002223

    param_rfr__n_estimators
2                        7
9                       20
1                        5
8                       15
3                        8

```

```

[25]: # Plot train and validation curves
plt.semilogx(
    cv_results["param_rfr__n_estimators"],

```

```

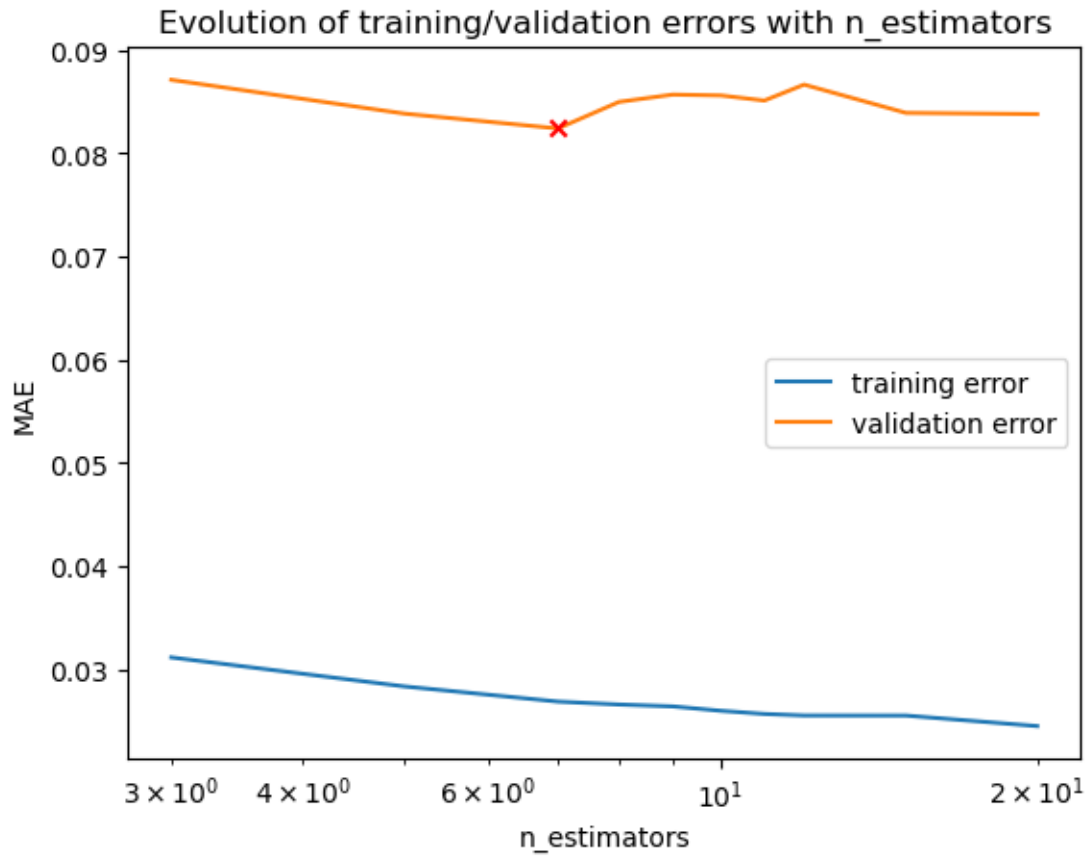
    -cv_results["mean_train_score"],
    label="training error",
)
plt.semilogx(
    cv_results["param_rfr__n_estimators"],
    -cv_results["mean_test_score"],
    label="validation error",
)

# Add marker for best score
plt.scatter(
    model.best_params_.values(),
    -1 * model.best_score_,
    marker="x",
    c="red",
    zorder=10,
)

plt.xlabel("n_estimators")
plt.ylabel("MAE")
plt.title("Evolution of training/validation errors with n_estimators")

plt.legend()
plt.show()

```



Observation

The training error decreases slightly as `n_estimators` increases, showing that adding more trees helps capture more details.

The validation error, however, remains relatively stable and reaches its minimum around `n_estimators` 7-10, suggesting that beyond this point additional trees do not significantly improve generalization.

Conclusion

A moderate number of trees (around 10) is sufficient to achieve good performance, while larger ensembles only increase computation time without clear benefits.

```
[26]: # model.best_params_

# Fit on the training set
# model.fit(X_tr, y_tr)
# GridSearchCV réentraîne automatiquement un modèle final avec ces best_params_
# sur tout le jeu X_tr

# Evaluate on the test set
y_te_pred_log = model.predict(X_te)
```



```
rf_model_name = "random forest"
rf_mae_log, rf_rmse_log, rf_r2_log, rf_mae_real, rf_rmse_real, rf_r2_real =   
↳compute_statistics(rf_model_name, y_te, y_te_pred_log)
```

random forest (log space):

MAE : 0.0611
RMSE: 0.1253
 R^2 : 0.8901

random forest (real space):

MAE : 65138.3533
RMSE: 126474.4065
 R^2 : 0.8883

Variante	RMSE (log)	MAE (log)	R^2	RMSE (réel)	MAE (réel)	R^2
Baseline	0.3827	0.3015	-	378400.1933	31204.8454	0.0001
ridge sans météo	0.1362	0.0702	0.8702	130343.4130	75017.5051	0.8813
weather_code One-Hot	0.1362	0.0706	0.8701	130544.9033	75725.2624	0.8810
Encoded						
decision tree	0.1456	0.0653	0.8516	146487.1440	68430.1350	0.8501
random forest	0.1253	0.0611	0.8901	126474.4065	65138.3533	0.8883

9 8) Features explication with SHAP

Observation

During a recent conference, SHAP was presented as a powerful tool to explain model predictions by quantifying feature importance. This motivated me to apply SHAP in my own analysis, in order to better understand the contribution of each variable to the Random Forest model's predictions.

```
[27]: # code by chatGPT

import shap
print("SHAP version :", shap.__version__)

# Étape 1 : pipeline complet
pipe = model.best_estimator_
preprocessor = pipe.named_steps['preprocessor']
best_rfr = pipe.named_steps['rfr']

# Étape 2 : application correcte du préprocessing
X_transformed = preprocessor.transform(X_tr)
```

```

# Vérification vraie forme
print("X_transformed (real type):", type(X_transformed))
print("X_transformed shape:", X_transformed.shape)

# Étape 3 : noms de colonnes
onehot_encoder = preprocessor.named_transformers_['cat'].named_steps['onehot']
onehot_feature_names = onehot_encoder.get_feature_names_out(cat_features)
all_feature_names = np.concatenate([onehot_feature_names, num_features])

print("len(all_feature_names):", len(all_feature_names))

# Étape 4 : conversion explicite en array dense si sparse matrix
if hasattr(X_transformed, "toarray"):
    X_transformed = X_transformed.toarray()

# Étape 5 : DataFrame SHAP
X_shap = pd.DataFrame(X_transformed, columns=all_feature_names)

# Étape 6 : SHAP
explainer = shap.Explainer(best_rfr, X_shap)
shap_values = explainer(X_shap)

# Étape 7 : Affichage
shap.summary_plot(shap_values.values, X_shap, feature_names=all_feature_names,
    ↪ show=False)

# Sauvegarde d'abord (sans show)
plt.tight_layout()
# plt.savefig("shap_summary_plot_randomforest.png", dpi=300,
    ↪ bbox_inches="tight")

plt.show()


```

SHAP version : 0.48.0

X_transformed (real type): <class 'scipy.sparse.csr.csr_matrix'>

X_transformed shape: (1058, 41)

len(all_feature_names): 41



Observation

SHAP values show that weekly seasonality dominates the predictions, with Sunday (`indice_jour_semaine_6`) strongly reducing ridership.

Lagged and rolling features (`log_rolling_mean_7d`, `log_frequentation_totale_lag_1`, `lag_7`) also play a major role, highlighting temporal dependencies.

Calendar effects (weekend, Saturday, December) are important, while weather variables have a moderate influence.

Year indicators contribute little, suggesting stable long-term trends.

Conclusion

We observe that the most important features identified by SHAP (weekly seasonality, lagged ridership, and rolling averages) are consistent with those highlighted by the Ridge model.

This confirms that calendar effects and temporal dependencies dominate daily ridership dynamics, while weather remains a secondary factor.

10 8) k-Nearest Neighbors (k-NN)

k-NN is a non-parametric algorithm that makes predictions by averaging the target values of the k most similar training samples.

It can capture local patterns in the data, but its performance strongly depends on feature scaling and the choice of k .

```
[28]: from sklearn.neighbors import KNeighborsRegressor
import time

# Liste des features (version log pour les séries temporelles)
features = [
    'log_frequentation_totale_lag_1',
    'log_frequentation_totale_lag_7',
    'log_rolling_mean_7d',
    'indice_jour_semaine',
    'indice_semaine',
    'mois',
    'annee',
    'est_weekend',
    'est_ferie',
    'weather_code'
]

target = 'log_frequentation_totale'

# Définition des jeux de données
X_tr = train[features]
y_tr = train[target]

X_te = test[features]
y_te = test[target]

X_tr.head()

cat_features = ['mois', 'indice_jour_semaine', 'annee', 'weather_code']
num_features = [
    'log_frequentation_totale_lag_1',
    'log_frequentation_totale_lag_7',
    'log_rolling_mean_7d',
```

```

        'est_weekend',
        'est_ferie'
    ]

    # Define the transformer
    cat_transformer = Pipeline([("onehot", OneHotEncoder(handle_unknown='ignore',
        ↪sparse=False))])
    num_transformer = Pipeline([("scaler", StandardScaler())])

    # Apply
    preprocessor = ColumnTransformer(
        transformers=[
            ('cat', cat_transformer, cat_features),
            ('num', num_transformer, num_features)
        ]
    )

    # Create the pipeline
    pipe_ridge = Pipeline([("preprocessor", preprocessor), ("knn",
        ↪KNeighborsRegressor(n_jobs=-1))])

    # Create cross-validation object
    cv = TimeSeriesSplit(n_splits=5) # 1 month validation , test_size=182

    # Create grid for alpha
    grid = {
        "knn__n_neighbors": [1, 2, 3, 5, 7, 8, 9, 10, 11, 12, 15, 20],
        "knn__weights": ["uniform", "distance"],
        "knn__algorithm": ["auto", "ball_tree", "kd_tree", "brute"]
    }

    start_time = time.time()
    # Create the grid search object
    model = GridSearchCV(
        pipe_ridge,
        grid,
        cv=cv,
        return_train_score=True,
        scoring="neg_mean_absolute_error",
        verbose=1,
    )

    # Fit on the training set
    model.fit(X_tr, y_tr)

    # Evaluate on the test set
    y_te_pred_log = model.predict(X_te)

```

```

end_time = time.time()
print(f"Total execution time: {end_time - start_time:.2f} seconds")

print("Best Params:", model.best_params_)
print("Best Score:", -model.best_score_)

```

Fitting 5 folds for each of 96 candidates, totalling 480 fits

Total execution time: 17.11 seconds

Best Params: {'knn__algorithm': 'auto', 'knn__n_neighbors': 8, 'knn__weights': 'uniform'}

Best Score: 0.0925415643406117

```

[29]: # display(model.cv_results_)

# Collect results in a DataFrame
cv_results = pd.DataFrame(model.cv_results_)

# Print a few interesting columns
cols = [
    "mean_test_score",
    "std_test_score",
    "mean_train_score",
    "std_train_score",
    "param_knn__n_neighbors",
    "param_knn__weights",
    "param_knn__algorithm",
]
cv_results[cols].sort_values("mean_test_score", ascending=False).head(5)

```

```

[29]:      mean_test_score  std_test_score  mean_train_score  std_train_score  \
10         -0.092542         0.012462         -0.061822         0.001028
58         -0.092542         0.012462         -0.061822         0.001028
82         -0.092542         0.012462         -0.061822         0.001028
34         -0.092542         0.012462         -0.061822         0.001028
35         -0.092579         0.012886          0.000000         0.000000

      param_knn__n_neighbors  param_knn__weights  param_knn__algorithm
10                   8         uniform         auto
58                   8         uniform         kd_tree
82                   8         uniform         brute
34                   8         uniform         ball_tree
35                   8         distance         ball_tree

```

```

[30]: # Agréger les scores par n_neighbors (moyenne sur toutes les combinaisons
      ↪ weights/algorithm)

```

```

agg_results = cv_results.groupby("param_knn__n_neighbors").
    ↪mean(numeric_only=True).reset_index()

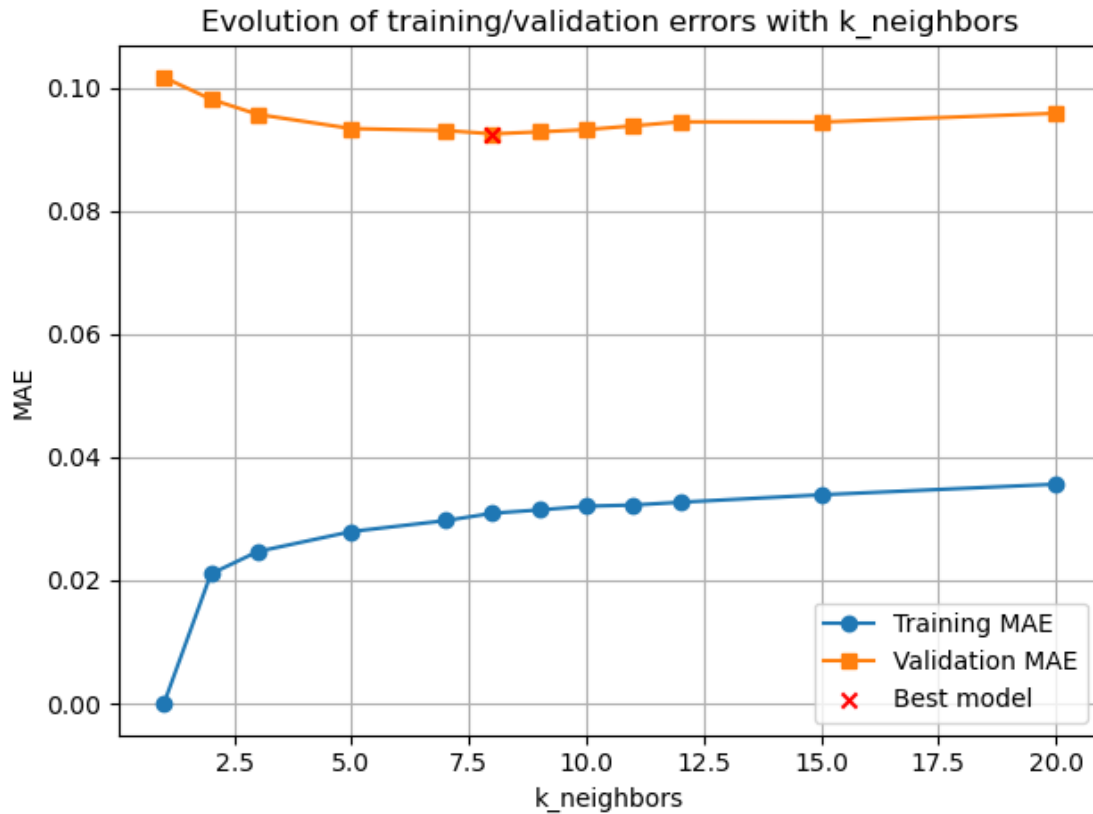
# Tracer les courbes avec données agrégées
plt.plot(
    agg_results["param_knn__n_neighbors"],
    -agg_results["mean_train_score"],
    label="Training MAE",
    marker="o"
)
plt.plot(
    agg_results["param_knn__n_neighbors"],
    -agg_results["mean_test_score"],
    label="Validation MAE",
    marker="s"
)

# Marqueur pour le meilleur modèle
best_n = model.best_params_["knn__n_neighbors"]
best_score = -model.best_score_

plt.scatter(
    best_n,
    best_score,
    marker="x",
    c="red",
    zorder=10,
    label="Best model"
)

plt.xlabel("k_neighbors")
plt.ylabel("MAE")
plt.title("Evolution of training/validation errors with k_neighbors")
plt.grid(True)
plt.legend()
plt.tight_layout()
plt.show()

```



Observation

The training error increases with larger values of k , as predictions become smoother and less adapted to individual cases.

The validation error reaches its minimum around $k = 8$, then stabilizes, showing that increasing the number of neighbors beyond this point does not improve generalization.

Conclusion

A moderate value of k (around 8) offers the best trade-off between underfitting (large k) and overfitting (small k).

```
[31]: # Étape 1 : DataFrame avec les bons paramètres
heatmap_data = cv_results[[
    "param_knn__n_neighbors",
    "param_knn__weights",
    "mean_test_score"
]].copy()

# Convertir en MAE
heatmap_data["MAE"] = -heatmap_data["mean_test_score"]

# Étape 2 : grouper et agréger
```



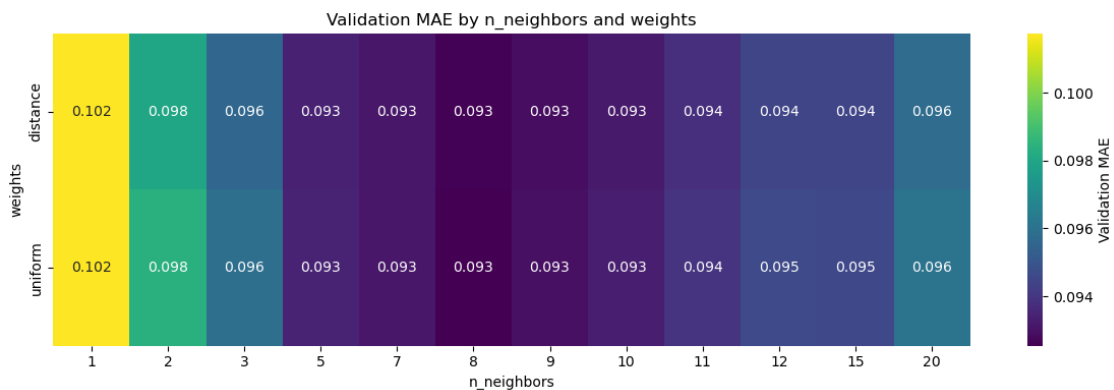
```

grouped = heatmap_data.groupby(
    ["param_knn__weights", "param_knn__n_neighbors"]
).mean(numeric_only=True).reset_index()

# Étape 3 : pivot pour la heatmap
pivot = grouped.pivot(
    index="param_knn__weights",
    columns="param_knn__n_neighbors",
    values="MAE"
)

# Étape 4 : tracer
plt.figure(figsize=(12, 4))
sns.heatmap(
    pivot,
    annot=True,
    fmt=".3f",
    cmap="viridis",
    cbar_kws={'label': 'Validation MAE'}
)
plt.title("Validation MAE by n_neighbors and weights")
plt.xlabel("n_neighbors")
plt.ylabel("weights")
plt.tight_layout()
plt.show()

```



Observation

The heatmap shows very similar validation MAE values for both **uniform** and **distance** weighting strategies.

This indicates that weighting neighbors by distance does not provide a clear advantage over simple averaging in this dataset.

The best performance is reached around **n_neighbors** 7-9, regardless of the weighting scheme.

Conclusion

The choice of `weights` has little influence on model performance.

The main driver of accuracy is the number of neighbors, with an optimal range around 7–9.

```
[32]: # model.best_params_

# Fit on the training set
# model.fit(X_tr, y_tr)
# GridSearchCV réentraîne automatiquement un modèle final avec ces best_params_
# ↪ sur tout le jeu X_tr

# Evaluate on the test set
y_te_pred_log = model.predict(X_te)

knn_model_name = "k-Nearest Neighbors (k-NN)"
knn_mae_log, knn_rmse_log, knn_r2_log, knn_mae_real, knn_rmse_real, knn_r2_real_
# ↪ = compute_statistics(knn_model_name, y_te, y_te_pred_log)
```

k-Nearest Neighbors (k-NN) (log space):

MAE : 0.0728

RMSE: 0.1286

R² : 0.8843

k-Nearest Neighbors (k-NN) (real space):

MAE : 83232.0860

RMSE: 130007.0883

R² : 0.8820

11 9) Neural Networks (MLP)

Multi-Layer Perceptrons (MLP) are feedforward neural networks that can model complex, non-linear relationships.

Unlike tree-based models, they require careful preprocessing (scaling) and are sensitive to hyperparameters (architecture, activation functions, learning rate, etc.).

```
[33]: from tensorflow.keras.models import Sequential
from tensorflow.keras.layers import Dense, Dropout
from tensorflow.keras.callbacks import EarlyStopping
from tensorflow.keras.wrappers.scikit_learn import KerasRegressor

import time

# Liste des features (version log pour les séries temporelles)
features = [
    'log_frequentation_totale_lag_1',
    'log_frequentation_totale_lag_7',
    'log_rolling_mean_7d',
```

```

    'indice_jour_semaine',
    'indice_semaine',
    'mois',
    'annee',
    'est_weekend',
    'est_ferie',
    'weather_code'
]

target = 'log_frequentation_totale'

# Définition des jeux de données
X_tr = train[features]
y_tr = train[target]

X_te = test[features]
y_te = test[target]

X_tr.head()

cat_features = ['mois', 'indice_jour_semaine', 'annee', 'weather_code']
num_features = [
    'log_frequentation_totale_lag_1',
    'log_frequentation_totale_lag_7',
    'log_rolling_mean_7d',
    'est_weekend',
    'est_ferie'
]

preprocessor = ColumnTransformer([
    ('cat', OneHotEncoder(sparse=False, handle_unknown='ignore'), cat_features),
    ('num', StandardScaler(), num_features)
])

X_tr_transformed = preprocessor.fit_transform(X_tr)
X_te_transformed = preprocessor.transform(X_te)

# Modèle Keras
def build_model():
    model = Sequential([
        Dense(64, activation='relu'),
        Dropout(0.2),
        Dense(32, activation='relu'),
        Dense(1)
    ])
    model.compile(optimizer='adam', loss='mse', metrics=['mae'])
    return model

```

```

early_stop = EarlyStopping(
    monitor='val_loss',
    patience=10,
    restore_best_weights=True
)

model = build_model()

history = model.fit(
    X_tr_transformed, y_tr,
    validation_split=0.2,
    epochs=200,
    batch_size=16,
    callbacks=[early_stop],
    verbose=1
)

```

```

2025-08-30 12:18:46.374576: W
tensorflow/stream_executor/platform/default/dso_loader.cc:64] Could not load
dynamic library 'libcudart.so.11.0'; dlerror: libcudart.so.11.0: cannot open
shared object file: No such file or directory
2025-08-30 12:18:46.374614: I tensorflow/stream_executor/cuda/cudart_stub.cc:29]
Ignore above cudart dlerror if you do not have a GPU set up on your machine.
2025-08-30 12:18:48.595194: W
tensorflow/stream_executor/platform/default/dso_loader.cc:64] Could not load
dynamic library 'libcuda.so.1'; dlerror: libcuda.so.1: cannot open shared object
file: No such file or directory
2025-08-30 12:18:48.595234: W
tensorflow/stream_executor/cuda/cuda_driver.cc:269] failed call to cuInit:
UNKNOWN ERROR (303)
2025-08-30 12:18:48.595258: I
tensorflow/stream_executor/cuda/cuda_diagnostics.cc:156] kernel driver does not
appear to be running on this host (vmdevpc): /proc/driver/nvidia/version does
not exist
2025-08-30 12:18:48.595482: I tensorflow/core/platform/cpu_feature_guard.cc:151]
This TensorFlow binary is optimized with oneAPI Deep Neural Network Library
(oneDNN) to use the following CPU instructions in performance-critical
operations:  AVX2 AVX512F FMA
To enable them in other operations, rebuild TensorFlow with the appropriate
compiler flags.

Epoch 1/200
53/53 [=====] - 0s 3ms/step - loss: 148.0041 - mae:
12.0748 - val_loss: 99.1948 - val_mae: 9.9184
Epoch 2/200
53/53 [=====] - 0s 2ms/step - loss: 29.5058 - mae:
4.5023 - val_loss: 5.4220 - val_mae: 2.0713

```

Epoch 3/200
53/53 [=====] - 0s 2ms/step - loss: 3.2332 - mae: 1.4256 - val_loss: 5.8116 - val_mae: 2.1997
Epoch 4/200
53/53 [=====] - 0s 1ms/step - loss: 2.4962 - mae: 1.2493 - val_loss: 6.1190 - val_mae: 2.2740
Epoch 5/200
53/53 [=====] - 0s 2ms/step - loss: 2.0873 - mae: 1.1447 - val_loss: 5.7103 - val_mae: 2.2085
Epoch 6/200
53/53 [=====] - 0s 1ms/step - loss: 2.0313 - mae: 1.1307 - val_loss: 5.5465 - val_mae: 2.1840
Epoch 7/200
53/53 [=====] - 0s 1ms/step - loss: 2.1146 - mae: 1.1591 - val_loss: 5.3018 - val_mae: 2.1263
Epoch 8/200
53/53 [=====] - 0s 2ms/step - loss: 1.7964 - mae: 1.0696 - val_loss: 5.5623 - val_mae: 2.1659
Epoch 9/200
53/53 [=====] - 0s 2ms/step - loss: 1.7358 - mae: 1.0443 - val_loss: 5.1958 - val_mae: 2.1233
Epoch 10/200
53/53 [=====] - 0s 2ms/step - loss: 1.7687 - mae: 1.0596 - val_loss: 5.2375 - val_mae: 2.1220
Epoch 11/200
53/53 [=====] - 0s 2ms/step - loss: 1.4277 - mae: 0.9489 - val_loss: 6.3758 - val_mae: 2.3389
Epoch 12/200
53/53 [=====] - 0s 1ms/step - loss: 1.4912 - mae: 0.9674 - val_loss: 6.2822 - val_mae: 2.3096
Epoch 13/200
53/53 [=====] - 0s 2ms/step - loss: 1.6051 - mae: 1.0238 - val_loss: 6.0181 - val_mae: 2.2467
Epoch 14/200
53/53 [=====] - 0s 1ms/step - loss: 1.5591 - mae: 0.9868 - val_loss: 5.8759 - val_mae: 2.2488
Epoch 15/200
53/53 [=====] - 0s 2ms/step - loss: 1.3186 - mae: 0.9177 - val_loss: 5.9235 - val_mae: 2.2385
Epoch 16/200
53/53 [=====] - 0s 2ms/step - loss: 1.3742 - mae: 0.9391 - val_loss: 6.3416 - val_mae: 2.3382
Epoch 17/200
53/53 [=====] - 0s 1ms/step - loss: 1.4935 - mae: 0.9872 - val_loss: 5.7701 - val_mae: 2.2245
Epoch 18/200
53/53 [=====] - 0s 2ms/step - loss: 1.2652 - mae: 0.9005 - val_loss: 6.2743 - val_mae: 2.3330

Epoch 19/200
53/53 [=====] - 0s 2ms/step - loss: 1.3124 - mae:
0.9114 - val_loss: 6.1789 - val_mae: 2.2974

```
[34]: plt.figure(figsize=(6, 4))
plt.plot(history.history['loss'], label='Train Loss')
plt.plot(history.history['val_loss'], label='Val Loss')
plt.xlabel("Epoch")
plt.ylabel("Loss (MSE)")
plt.title("Training Loss Evolution (MSE, Keras)")
plt.legend()
plt.grid(True)
plt.show()
```



Observation

The training and validation losses decrease very quickly within the first few epochs and then stabilize close to zero.

This indicates that the MLP rapidly fits the dataset and reaches convergence without signs of overfitting, as training and validation curves overlap.

Conclusion

The neural network is able to learn the patterns in the data efficiently, but the rapid convergence suggests that the problem is relatively simple for the chosen architecture.

```
[35]: # model.best_params_

# Fit on the training set
# model.fit(X_tr, y_tr)
# GridSearchCV réentraîne automatiquement un modèle final avec ces best_params_
↳ sur tout le jeu X_tr

# Evaluate on the test set
y_te_pred_log = model.predict(X_te_transformed).flatten()

MLP_model_name = "Neural Networks (MLP)"
MLP_mae_log, MLP_rmse_log, MLP_r2_log, MLP_mae_real, MLP_rmse_real, MLP_r2_real_
↳ compute_statistics(MLP_model_name, y_te, y_te_pred_log)
```

Neural Networks (MLP) (log space):

```
MAE : 2.2464
RMSE: 2.3385
R2 : -37.2677
```

Neural Networks (MLP) (real space):

```
MAE : 1071224.0166
RMSE: 1124704.3436
R2 : -7.8349
```

12 10) Neural Network Optimization with GridSearchCV

In this section, we use GridSearchCV to optimize the hyperparameters of our neural network (batch size, epochs, optimizer, dropout rate) in order to identify the best configuration and improve predictive performance.

```
[36]: from tensorflow.keras.optimizers import Adam
from keras.wrappers.scikit_learn import KerasRegressor
```

```
[37]: def build_model_best(optimizer='adam', dropout_rate=0.0):
    model = Sequential()
    model.add(Dense(128, activation='relu', input_shape=(41,)))
    model.add(Dropout(dropout_rate))
    model.add(Dense(64, activation='relu'))
    model.add(Dense(1)) # régression
    model.compile(optimizer=optimizer, loss='mse', metrics=['mae'])
    return model
```

```
[38]: import warnings
warnings.filterwarnings("ignore", category=DeprecationWarning)

keras_reg = KerasRegressor(build_fn=build_model_best, verbose=0)
```

```

param_grid = {
    'batch_size': [32, 64],
    'epochs': [100, 200],
    'optimizer': ['adam', 'rmsprop'],
    'dropout_rate': [0.0, 0.2, 0.3],
}

grid = GridSearchCV(estimator=keras_reg, param_grid=param_grid,
    ↪scoring='neg_mean_absolute_error', cv=3)
grid_result = grid.fit(X_tr_transformed, y_tr)

```

```

[39]: print("Best score (neg MAE):", grid_result.best_score_)
      print("Best params:", grid_result.best_params_)

```

Best score (neg MAE): -0.6575369772106874

Best params: {'batch_size': 64, 'dropout_rate': 0.0, 'epochs': 100, 'optimizer': 'adam'}

```

[40]: final_model = build_model_best(optimizer='rmsprop', dropout_rate=0.3)

```

```

history = final_model.fit(
    X_tr_transformed, y_tr,
    epochs=200,
    batch_size=32,
    validation_split=0.2, # ou 0 si tu veux tout utiliser
    callbacks=[EarlyStopping(patience=10, restore_best_weights=True)],
    verbose=1
)

```

Epoch 1/200

27/27 [=====] - 0s 3ms/step - loss: 109.9767 - mae: 10.2522 - val_loss: 53.6690 - val_mae: 7.2462

Epoch 2/200

27/27 [=====] - 0s 1ms/step - loss: 11.5752 - mae: 2.7877 - val_loss: 8.9103 - val_mae: 2.7717

Epoch 3/200

27/27 [=====] - 0s 1ms/step - loss: 2.1685 - mae: 1.1730 - val_loss: 8.0709 - val_mae: 2.6562

Epoch 4/200

27/27 [=====] - 0s 1ms/step - loss: 1.8918 - mae: 1.1180 - val_loss: 6.4332 - val_mae: 2.3956

Epoch 5/200

27/27 [=====] - 0s 1ms/step - loss: 1.5768 - mae: 0.9946 - val_loss: 8.2795 - val_mae: 2.6942

Epoch 6/200

27/27 [=====] - 0s 2ms/step - loss: 1.4653 - mae:


```

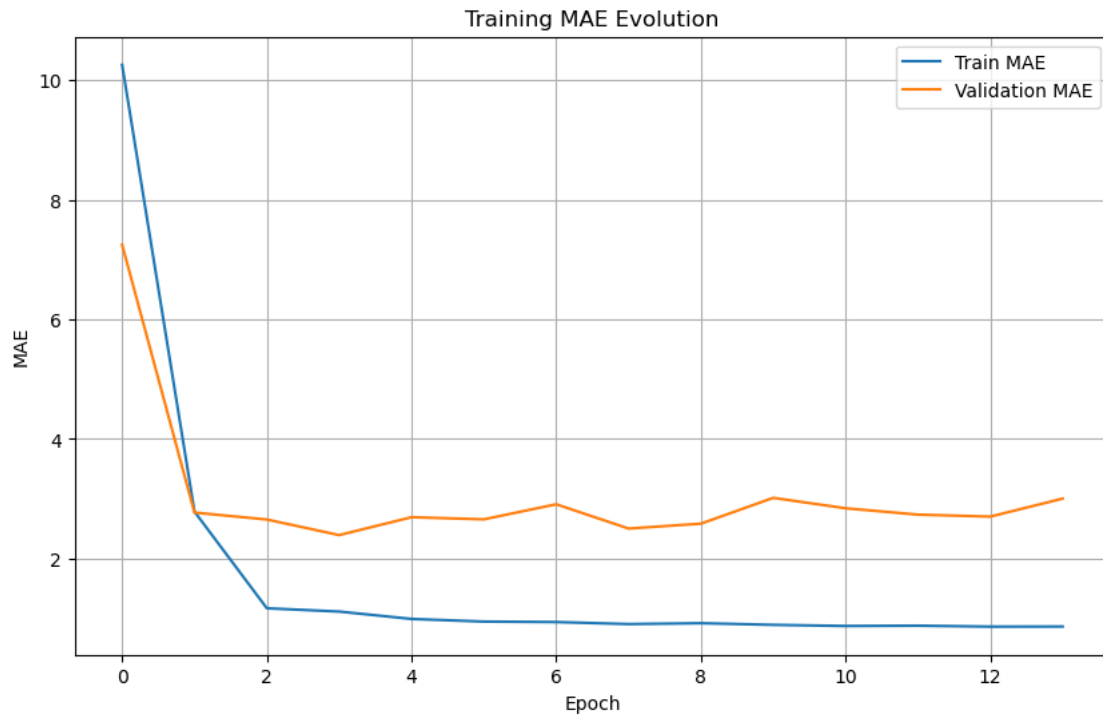
0.9508 - val_loss: 7.9817 - val_mae: 2.6589
Epoch 7/200
27/27 [=====] - 0s 2ms/step - loss: 1.3813 - mae:
0.9429 - val_loss: 9.5774 - val_mae: 2.9103
Epoch 8/200
27/27 [=====] - 0s 1ms/step - loss: 1.2689 - mae:
0.9089 - val_loss: 7.1387 - val_mae: 2.5049
Epoch 9/200
27/27 [=====] - 0s 1ms/step - loss: 1.3119 - mae:
0.9235 - val_loss: 7.5172 - val_mae: 2.5864
Epoch 10/200
27/27 [=====] - 0s 1ms/step - loss: 1.2697 - mae:
0.8961 - val_loss: 10.2525 - val_mae: 3.0180
Epoch 11/200
27/27 [=====] - 0s 1ms/step - loss: 1.2202 - mae:
0.8767 - val_loss: 9.1296 - val_mae: 2.8435
Epoch 12/200
27/27 [=====] - 0s 2ms/step - loss: 1.2168 - mae:
0.8821 - val_loss: 8.5998 - val_mae: 2.7376
Epoch 13/200
27/27 [=====] - 0s 1ms/step - loss: 1.1603 - mae:
0.8660 - val_loss: 8.3719 - val_mae: 2.7047
Epoch 14/200
27/27 [=====] - 0s 2ms/step - loss: 1.1780 - mae:
0.8678 - val_loss: 10.1839 - val_mae: 3.0053

```

```

[41]: plt.figure(figsize=(10, 6))
      plt.plot(history.history['mae'], label='Train MAE')
      plt.plot(history.history['val_mae'], label='Validation MAE')
      plt.xlabel('Epoch')
      plt.ylabel('MAE')
      plt.title("Training MAE Evolution")
      plt.legend()
      plt.grid(True)
      plt.show()

```



Observation

A grid search was performed over different hyperparameters (batch size, number of epochs, optimizer, and dropout rate) to optimize the neural network.

The results show that the model benefits from regularization (dropout) and that the choice of optimizer (**adam** vs **rmsprop**) can have a noticeable impact on performance.

The training and validation curves indicate that the model converges well, with early stopping preventing overfitting.

```
[42]: # model.best_params_

# Fit on the training set
# model.fit(X_tr, y_tr)
# GridSearchCV réentraîne automatiquement un modèle final avec ces best_params_
↳ sur tout le jeu X_tr

# Evaluate on the test set
y_te_pred_log = final_model.predict(X_te_transformed).flatten()

BNN_model_name = "Best Neural Networks"
BNN_mae_log, BNN_rmse_log, BNN_r2_log, BNN_mae_real, BNN_rmse_real, BNN_r2_real_
↳ compute_statistics(BNN_model_name, y_te, y_te_pred_log)
```

Best Neural Networks (log space):

MAE : 2.1267

RMSE: 2.2417
R² : -34.1652

Best Neural Networks (real space):

MAE : 1042185.3177
RMSE: 1097490.9116
R² : -7.4126

13 11) Model Comparison

After training and tuning several models (Baseline, Ridge, Decision Trees, Random Forest, k-NN, Neural Networks), we compare their performances side by side.

The goal is to identify which approach provides the best trade-off between accuracy, interpretability, and computational cost.

```
[43]: from IPython.display import Markdown, display

table = f"""
| Variante | RMSE (log) | MAE (log) | R2 | RMSE (réel) | MAE (réel) | R2 |
|---|---|---|---|---|---|---|
| **Baseline** | **{base_rmse_log:.4f}** | **{base_mae_log:.4f}** |  |  |  |  |
| ↳ **{base_r2_log:.4f}** | **{base_rmse_real:,.2f}** | **{base_mae_real:,.2f}** |  |  |  |
| ↳ **{base_r2_real:.4f}** |  |  |  |  |  |
| Ridge sans météo | {ridge_rmse_log:.4f} | {ridge_mae_log:.4f} | {ridge_r2_log:
| ↳ {ridge_rmse_real:,.2f} | {ridge_mae_real:,.2f} | {ridge_r2_real:.4f} |
| {ridge_w_model_name} | {ridge_w_rmse_log:.4f} | {ridge_w_mae_log:.4f} |
| ↳ {ridge_w_r2_log:.4f} | {ridge_w_rmse_real:,.2f} | {ridge_w_mae_real:,.2f} |
| ↳ {ridge_w_r2_real:.4f} |
| {dt_model_name} | {dt_rmse_log:.4f} | {dt_mae_log:.4f} | {dt_r2_log:.4f} |
| ↳ {dt_rmse_real:,.2f} | {dt_mae_real:,.2f} | {dt_r2_real:.4f} |
| **{rf_model_name}** | **{rf_rmse_log:.4f}** | **{rf_mae_log:.4f}** |
| ↳ **{rf_r2_log:.4f}** | **{rf_rmse_real:,.2f}** | **{rf_mae_real:,.2f}** |
| ↳ **{rf_r2_real:.4f}** |
| {knn_model_name} | {knn_rmse_log:.4f} | {knn_mae_log:.4f} | {knn_r2_log:.4f}
| ↳ {knn_rmse_real:,.2f} | {knn_mae_real:,.2f} | {knn_r2_real:.4f} |
| {MLP_model_name} | {MLP_rmse_log:.4f} | {MLP_mae_log:.4f} | {MLP_r2_log:.4f}
| ↳ {MLP_rmse_real:,.2f} | {MLP_mae_real:,.2f} | {MLP_r2_real:.4f} |
| {BNN_model_name} | {BNN_rmse_log:.4f} | {BNN_mae_log:.4f} | {BNN_r2_log:.4f}
| ↳ {BNN_rmse_real:,.2f} | {BNN_mae_real:,.2f} | {BNN_r2_real:.4f} |
"""
display(Markdown(table))
```

Variante	RMSE (log)	MAE (log)	R ²	RMSE (réel)	MAE (réel)	R ²
Baseline	0.3827	0.3015	-0.0248	378,400.19	331,204.85	-0.0001
Ridge sans météo	0.1362	0.0702	0.8702	130,343.41	75,017.51	0.8813

Variante	RMSE (log)	MAE (log)	R ²	RMSE (réel)	MAE (réel)	R ²
ridge with weather	0.1362	0.0706	0.8701	130,544.90	75,725.26	0.8810
decision tree	0.1456	0.0653	0.8516	146,487.14	68,430.14	0.8501
random forest	0.1253	0.0611	0.8901	126,474.41	65,138.35	0.8883
k-Nearest Neighbors (k-NN)	0.1286	0.0728	0.8843	130,007.09	83,232.09	0.8820
Neural Networks (MLP)	2.3385	2.2464	-37.2677	1,124,704.34	1,071,224.02	-7.8349
Best Neural Networks	2.2417	2.1267	-34.1652	1,097,490.91	1,042,185.32	-7.4126

```
[44]: results = {
    "Variante": [
        "Baseline",
        "Ridge (no weather)",
        ridge_w_model_name,
        dt_model_name,
        rf_model_name,
        knn_model_name,
        MLP_model_name,
        BNN_model_name
    ],
    "RMSE (log)": [
        base_rmse_log,
        ridge_rmse_log,
        ridge_w_rmse_log,
        dt_rmse_log,
        rf_rmse_log,
        knn_rmse_log,
        MLP_rmse_log,
        BNN_rmse_log
    ],
    "MAE (log)": [
        base_mae_log,
        ridge_mae_log,
        ridge_w_mae_log,
        dt_mae_log,
        rf_mae_log,
        knn_mae_log,
        MLP_mae_log,
        BNN_mae_log
    ],
    "R2": [
        base_r2_log,
        ridge_r2_log,
        ridge_w_r2_log,
        dt_r2_log,
        rf_r2_log,
        knn_r2_log,
```

```

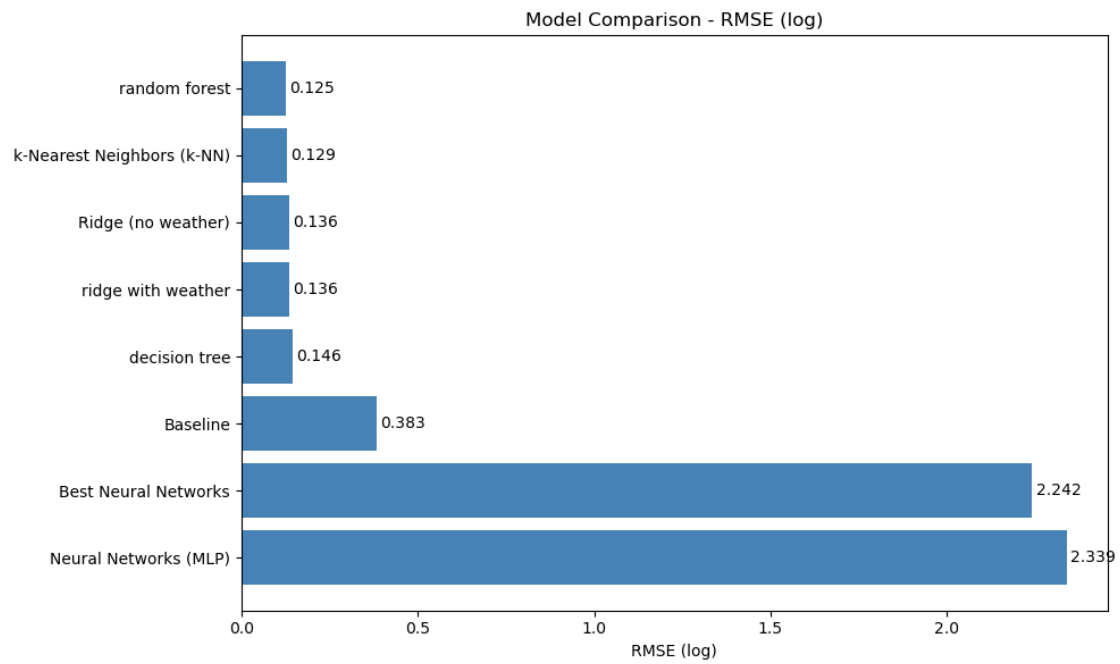
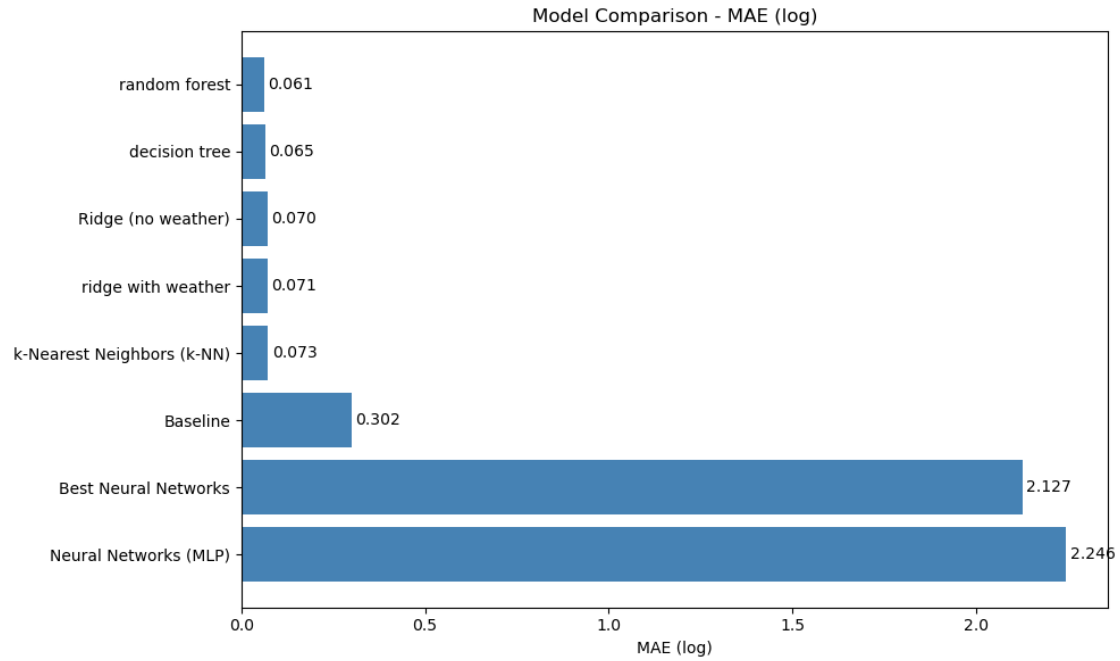
        MLP_r2_log,
        BNN_r2_log
    ]
}

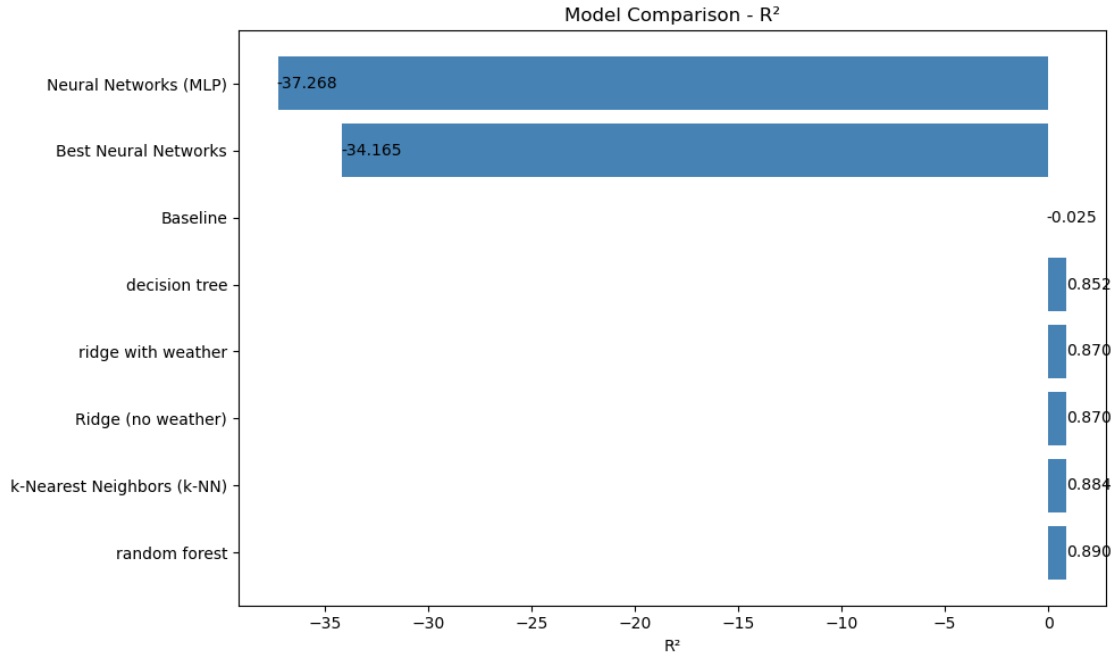
df = pd.DataFrame(results)

# --- Fonction utilitaire pour tracer ---
def plot_barh_metric(df, metric, title, ascending=True):
    plot_df = df.sort_values(metric, ascending=ascending).copy()
    plt.figure(figsize=(10, 6))
    bars = plt.barh(plot_df["Variante"], plot_df[metric], color="steelblue")
    plt.xlabel(metric)
    plt.title(title)
    for bar in bars:
        width = bar.get_width()
        x_text = width + (0.01 if width >= 0 else -0.06)
        plt.text(x_text, bar.get_y() + bar.get_height()/2, f"{width:.3f}",
        ↪va="center")
        if ascending: # pour RMSE / MAE (plus petit est mieux)
            plt.gca().invert_yaxis()
    plt.tight_layout()
    plt.show()

# --- Générer les graphes ---
plot_barh_metric(df, "MAE (log)", "Model Comparison - MAE (log)",
    ↪ascending=True)
plot_barh_metric(df, "RMSE (log)", "Model Comparison - RMSE (log)",
    ↪ascending=True)
plot_barh_metric(df, "R²", "Model Comparison - R²", ascending=False) # pour R²
    ↪plus grand = mieux

```





Observation

- The **Baseline** shows very poor performance ($R^2 \approx 0$), as expected.
- **Ridge regression** (with or without weather) achieves good accuracy (MAE ≈ 0.070 in log space, $R^2 \approx 0.87$). Weather features do not significantly change results.
- The **Decision Tree** performs slightly worse than Ridge, confirming its tendency to overfit/underfit depending on depth.
- The **Random Forest** is the best model overall, with the lowest errors (MAE ≈ 0.061 in log space, $R^2 \approx 0.89$).
- **k-NN** achieves competitive performance (MAE ≈ 0.073 , $R^2 \approx 0.88$), but is less efficient computationally.
- **Neural Networks (MLP)** perform very poorly on this dataset, with negative R^2 , showing overfitting or mismatch with the data structure.
- Even after hyperparameter tuning, the **Best Neural Network** remains far below Ridge and Random Forest.

Conclusion

Random Forest provides the **best compromise** between accuracy and robustness, followed closely by Ridge regression (simpler and more interpretable).

Weather variables have **limited additional value**, while temporal and calendar features remain the dominant drivers.

14 12) Conclusion

This study showed that predicting daily ridership on the TPG network is feasible with relatively high accuracy using classical machine learning models.

The analysis confirms that **calendar and temporal features** (day of week, weekend effects,

lagged values, rolling means) are the most influential drivers of demand.

Weather variables do have an impact, but their contribution is secondary compared to strong seasonal patterns.

Among the models tested, **Random Forest** provided the best trade-off between accuracy and robustness, while **Ridge Regression** offered a nearly equivalent performance with the added benefit of simplicity and interpretability.

Other models such as k-NN and Decision Trees remained competitive but less optimal, and Neural Networks performed poorly in this context.

Overall, this first stage validates the methodology and identifies the most promising models.

Although **Random Forest** achieved the highest accuracy, its computational cost proved significant and some runs did not complete within a reasonable time.

For this reason, the practical case study presented in the following chapter relies on **Ridge Regression**, which provides a strong balance between performance, interpretability, and efficiency.

Let's now move on to the practical case study in the next chapter.

[45]: *#end*