

5_ML_TPG_Stop

September 3, 2025

1 Machine Learning with Daily Boarding/Alighting Details by Stop

In this section, I aim to carry out a **practical case study**.
The objective is to answer a concrete question such as:

“How many boardings at stop XXX on line YYY on date ZZZ?”

This approach allows us to test a prediction model on a real and easily interpretable situation.

However, I was not able to implement a simple function that correctly recomputes the *lag* variables for the future period.

Therefore, I chose another approach: truncate the dataset up to **June 19, 2025**, estimate the number of boardings for **June 20, 2025**, and then compare this prediction with the actual value contained in the original dataset.

```
[1]: INSTALL_LIB = False
```

```
[2]: # libraries
      %matplotlib inline
      import numpy as np
      import pandas as pd
      import matplotlib.pyplot as plt
      import seaborn as sns
```

2 1) Loading the data

```
[3]: file_name = "TPG_daily_df.csv"

      TPG_meteo_all_df = pd.read_csv(file_name, sep=",", low_memory=False)
      TPG_meteo_all_df['date'] = pd.to_datetime(TPG_meteo_all_df['date'])
      display(TPG_meteo_all_df.info())
```

```
<class 'pandas.core.frame.DataFrame'>
RangeIndex: 4794267 entries, 0 to 4794266
Data columns (total 33 columns):
 #   Column          Dtype
---  -

```

```

0    date                                datetime64[ns]
1    ligne                               object
2    ligne_type_act                       object
3    horaire_type                         object
4    arret_code_long                      object
5    indice_semaine                      int64
6    indice_jour_semaine                  int64
7    nb_de_montees                        float64
8    nb_de_descentes                      float64
9    annee                                int64
10   mois                                 int64
11   log_montees                          float64
12   log_descentes                        float64
13   frequentation_totale                 float64
14   log_frequentation_totale             float64
15   delta_montees_descentes              float64
16   log_delta_montees_descentes          float64
17   type_vehicule                        object
18   weather_code                         int64
19   montees_lag_1                        float64
20   log_montees_lag_1                    float64
21   desc_lag_1                           float64
22   log_desc_lag_1                       float64
23   montees_lag_7                        float64
24   log_montees_lag_7                    float64
25   desc_lag_7                           float64
26   log_desc_lag_7                       float64
27   rolling_montees_7d                   float64
28   log_rolling_montees_7d               float64
29   rolling_desc_7d                      float64
30   log_rolling_desc_7d                  float64
31   est_weekend                          int64
32   est_ferie                            int64
dtypes: datetime64[ns](1), float64(20), int64(7), object(5)
memory usage: 1.2+ GB

```

None

```
[4]: TPG_meteo_all_df.head()
```

```

[4]:      date ligne ligne_type_act horaire_type arret_code_long \
0 2021-08-08      1    PRINCIPAL    DIMANCHE      31DC00
1 2021-08-09      1    PRINCIPAL    VACANCES      31DC00
2 2021-08-10      1    PRINCIPAL    VACANCES      31DC00
3 2021-08-11      1    PRINCIPAL    VACANCES      31DC00
4 2021-08-12      1    PRINCIPAL    VACANCES      31DC00

      indice_semaine  indice_jour_semaine  nb_de_montees  nb_de_descentes  annee \

```

0	31	7	88.06	70.72	2021
1	32	1	192.44	204.33	2021
2	32	2	205.16	224.59	2021
3	32	3	187.67	217.00	2021
4	32	4	214.80	217.32	2021

	montees_lag_7	log_montees_lag_7	desc_lag_7	log_desc_lag_7	\
0	...	86.84	4.475517	81.99	4.418720
1	...	149.95	5.016949	164.72	5.110300
2	...	199.67	5.301662	223.32	5.413074
3	...	235.60	5.466371	205.42	5.329913
4	...	219.47	5.395762	208.25	5.343530

	rolling_montees_7d	log_rolling_montees_7d	rolling_desc_7d	\
0	174.415714	5.167159	173.398571	
1	180.485714	5.201177	179.057143	
2	181.270000	5.205489	179.238571	
3	174.422857	5.167199	180.892857	
4	173.755714	5.163389	182.188571	

	log_rolling_desc_7d	est_weekend	est_ferie
0	5.161343	0	0
1	5.193274	0	0
2	5.194281	0	0
3	5.203418	0	0
4	5.210516	0	0

[5 rows x 33 columns]

2.0.1 Action :

To build the features for **June 20, 2025** and compare the final result,
I store the data for this date in a separate dataset.

```
[5]: TPG_20250620 = TPG_meteo_all_df[(TPG_meteo_all_df["date"] == "2025-06-20")]

#Debug, on affiche les données du 20/06/2025 pour l'arret 'RIVE00' sur la ligne
↪12
TPG_20250620[(TPG_20250620["arret_code_long"] == "RIVE00") &
↪(TPG_20250620["ligne"] == "12")].head(1)
```

```
[5]:          date ligne ligne_type_act horaire_type arret_code_long \
3748340 2025-06-20    12      PRINCIPAL      NORMAL      RIVE00

          indice_semaine  indice_jour_semaine  nb_de_montees  nb_de_descentes \
3748340                25                5        3320.14        3309.28
```

```

      annee ... montees_lag_7 log_montees_lag_7 desc_lag_7 \
3748340  2025 ...      2960.69      7.993515      3356.01

      log_desc_lag_7 rolling_montees_7d log_rolling_montees_7d \
3748340      8.118806      2491.258571      7.820945

      rolling_desc_7d log_rolling_desc_7d est_weekend est_ferie
3748340      2701.362857      7.901882      1      0

[1 rows x 33 columns]

```

2.0.2 Action :

We truncate the dataset at **June 19, 2025** in order to exclude the day of June 20.

```
[6]: cutoff = pd.to_datetime("2025-06-19")
TPG_meteo_all_df_cut = TPG_meteo_all_df[TPG_meteo_all_df["date"] <= cutoff]
```

```
[7]: # descriptive statistics
TPG_meteo_all_df_cut.describe().T
```

```
[7]:
```

	count	mean	std	min	\
indice_semaine	4784731.0	26.448799	15.265995	1.000000	
indice_jour_semaine	4784731.0	3.901318	1.949375	1.000000	
nb_de_montees	4784731.0	170.034720	429.162980	0.000000	
nb_de_descentes	4784731.0	170.099195	426.501738	0.000000	
annee	4784731.0	2023.032681	1.185014	2021.000000	
mois	4784731.0	6.510411	3.514600	1.000000	
log_montees	4784731.0	3.176362	2.185523	0.000000	
log_descentes	4784731.0	3.224178	2.158247	0.000000	
frequentation_totale	4784731.0	340.133915	803.560279	0.000000	
log_frequentation_totale	4784731.0	6.400540	4.004097	0.000000	
delta_montees_descentes	4784731.0	-0.064475	294.040933	-10918.140000	
log_delta_montees_descentes	4784731.0	-0.047816	1.684128	-9.086355	
weather_code	4784731.0	31.431903	28.224605	0.000000	
montees_lag_1	4784731.0	169.931735	428.985746	0.000000	
log_montees_lag_1	4784731.0	3.175562	2.185451	0.000000	
desc_lag_1	4784731.0	169.995231	426.327290	0.000000	
log_desc_lag_1	4784731.0	3.223368	2.158187	0.000000	
montees_lag_7	4784731.0	169.558615	428.331599	0.000000	
log_montees_lag_7	4784731.0	3.172503	2.185337	0.000000	
desc_lag_7	4784731.0	169.626037	425.665579	0.000000	
log_desc_lag_7	4784731.0	3.220279	2.158180	0.000000	
rolling_montees_7d	4784731.0	169.818729	412.354260	0.000000	
log_rolling_montees_7d	4784731.0	3.299882	2.103777	0.000000	
rolling_desc_7d	4784731.0	169.882173	409.400698	0.000000	
log_rolling_desc_7d	4784731.0	3.351139	2.070083	0.000000	

est_weekend	4784731.0	0.295849	0.456424	0.000000
est_ferie	4784731.0	0.000000	0.000000	0.000000

	25%	50%	75%	\
indice_semaine	13.000000	26.000000	40.000000	
indice_jour_semaine	2.000000	4.000000	6.000000	
nb_de_montees	2.840000	22.540000	137.550000	
nb_de_descentes	3.000000	24.790000	138.170000	
annee	2022.000000	2023.000000	2024.000000	
mois	3.000000	6.000000	10.000000	
log_montees	1.345472	3.158701	4.931231	
log_descentes	1.386294	3.249987	4.935696	
frequentation_totale	11.730000	64.150000	308.300000	
log_frequentation_totale	3.178054	6.185255	9.514157	
delta_montees_descentes	-29.860000	0.000000	27.260000	
log_delta_montees_descentes	-1.045502	0.000000	0.976823	
weather_code	3.000000	51.000000	61.000000	
montees_lag_1	2.830000	22.520000	137.430000	
log_montees_lag_1	1.342865	3.157851	4.930365	
desc_lag_1	3.000000	24.760000	138.050000	
log_desc_lag_1	1.386294	3.248823	4.934834	
montees_lag_7	2.820000	22.420000	137.010000	
log_montees_lag_7	1.340250	3.153590	4.927326	
desc_lag_7	3.000000	24.660000	137.610000	
log_desc_lag_7	1.386294	3.244933	4.931664	
rolling_montees_7d	3.614286	25.345714	143.357143	
log_rolling_montees_7d	1.529157	3.271306	4.972290	
rolling_desc_7d	3.861429	28.011429	145.435714	
log_rolling_desc_7d	1.581332	3.367690	4.986587	
est_weekend	0.000000	0.000000	1.000000	
est_ferie	0.000000	0.000000	0.000000	

	max
indice_semaine	52.000000
indice_jour_semaine	7.000000
nb_de_montees	9821.710000
nb_de_descentes	10944.990000
annee	2025.000000
mois	12.000000
log_montees	9.192452
log_descentes	9.300728
frequentation_totale	17919.470000
log_frequentation_totale	18.195097
delta_montees_descentes	8326.610000
log_delta_montees_descentes	8.885554
weather_code	75.000000
montees_lag_1	9821.710000

log_montees_lag_1	9.192452
desc_lag_1	10944.990000
log_desc_lag_1	9.300728
montees_lag_7	9821.710000
log_montees_lag_7	9.192452
desc_lag_7	10944.990000
log_desc_lag_7	9.300728
rolling_montees_7d	8453.642857
log_rolling_montees_7d	9.042471
rolling_desc_7d	9086.812857
log_rolling_desc_7d	9.114690
est_weekend	1.000000
est_ferie	0.000000

```
[8]: print(f"The data covers the date range from {TPG_meteo_all_df_cut['date'].min()} to {TPG_meteo_all_df_cut['date'].max()}")
```

The data covers the date range from 2021-08-08 00:00:00 to 2025-06-19 00:00:00

```
[9]: df = TPG_meteo_all_df_cut.copy()
```

3 2) Splitting the data

```
[10]: from datetime import datetime

# Définition des bornes
train_end = pd.to_datetime("2023-12-31")
val_end = pd.to_datetime("2024-06-30")

# Split temporel
train_df = df[df["date"] <= train_end]
val_df = df[(df["date"] > train_end) & (df["date"] <= val_end)]
test_df = df[df["date"] > val_end]
```

4 3) Target and features

4.0.1 Action :

We remove the raw values from the dataset and keep only their log-transformed versions.

```
[11]: drop_cols = [
    "date",          # index temporel
    "nb_de_montees", "nb_de_descentes",
    "log_montees",   # cible en log

    "frequentation_totale",
```

```

    "delta_montees_descentes",
    "montees_lag_1",
    "desc_lag_1",
    "montees_lag_7",
    "desc_lag_7",
    "rolling_montees_7d",
    "rolling_desc_7d"
]

```

```

[12]: X_tr = train_df.drop(columns=drop_cols)
      y_tr = train_df["log_montees"]

      X_val = val_df.drop(columns=drop_cols)
      y_val = val_df["log_montees"]

      X_te = test_df.drop(columns=drop_cols)
      y_te = test_df["log_montees"]

```

5 4) Model training

```

[13]: from sklearn.metrics import mean_squared_error, mean_absolute_error, r2_score

def compute_statistics(model_name, y_te, y_te_pred_log) :
    mae_log = mean_absolute_error(y_te, y_te_pred_log)
    rmse_log = np.sqrt(mean_squared_error(y_te, y_te_pred_log))
    r2_log = r2_score(y_te, y_te_pred_log)

    print(f"{model_name} (log space):")
    print(f"  MAE : {mae_log:.4f}")
    print(f"  RMSE: {rmse_log:.4f}")
    print(f"  R2  : {r2_log:.4f}")

    y_te_pred_real = np.exp(y_te_pred_log)
    y_te_real = np.exp(y_te)

    mae_real = mean_absolute_error(y_te_real, y_te_pred_real)
    rmse_real = np.sqrt(mean_squared_error(y_te_real, y_te_pred_real))
    r2_real = r2_score(y_te_real, y_te_pred_real)

    print("\n")
    print(f"{model_name} (real space):")
    print(f"  MAE : {mae_real:.4f}")
    print(f"  RMSE: {rmse_real:.4f}")
    print(f"  R2  : {r2_real:.4f}")

```

```
return mae_log, rmse_log, r2_log, mae_real, rmse_real, r2_real
```

```
[14]: from sklearn.pipeline import Pipeline
from sklearn.preprocessing import OneHotEncoder
from sklearn.compose import ColumnTransformer
from sklearn.preprocessing import StandardScaler

# Colonnes catégorielles (OneHot)
cat_features = [
    "ligne",          # identifiant ligne
    "ligne_type_act", # type de ligne
    "horaire_type",   # horaire (école, vacances, etc.)
    "arret_code_long", # arrêt
    "indice_semaine",
    "indice_jour_semaine",
    "annee",
    "mois",
    "type_vehicule",  # tram, bus, etc.
    "weather_code",   # météo catégorielle
    "est_weekend",
    "est_ferie"
]

# Colonnes numériques (scalées)
num_features = [
    "log_descentes",
    "log_frequentation_totale",
    "log_delta_montees_descentes",
    "log_montees_lag_1",
    "log_desc_lag_1",
    "log_montees_lag_7",
    "log_desc_lag_7",
    "log_rolling_montees_7d",
    "log_rolling_desc_7d"
]

# Define the transformer
cat_transformer = Pipeline([("onehot", OneHotEncoder(handle_unknown='ignore'))])
num_transformer = Pipeline([("scaler", StandardScaler())])

# Apply
preprocessor = ColumnTransformer(
    transformers=[
        ('cat', cat_transformer, cat_features),
        ('num', num_transformer, num_features)
    ]
)
```



```
)
```

```
[15]: from sklearn.linear_model import Ridge
      from sklearn.model_selection import GridSearchCV
      from sklearn.model_selection import TimeSeriesSplit
      import joblib
      import time

      # Create the pipeline
      pipe_ridge = Pipeline([("preprocessor", preprocessor), ("ridge", Ridge())])

      # Create cross-validation object
      cv = TimeSeriesSplit(n_splits=5)

      # Create grid for alpha
      grid = {"ridge__alpha": np.logspace(-4, 4, num=20)}

      # Create the grid search object
      model_ridge = GridSearchCV(
          pipe_ridge,
          grid,
          cv=cv,
          return_train_score=True,
          scoring="neg_mean_absolute_error",
          verbose=1,
      )

      # Fit on the training set
      start_time = time.time()
      model_ridge.fit(X_tr, y_tr)

      end_time = time.time()
      print(f"Total execution time: {end_time - start_time:.2f} seconds")

      print("Best Params:", model_ridge.best_params_)
      print("Best Score:", -model_ridge.best_score_)

      # save the model
      # joblib.dump(model_ridge, "ridge_pipeline_log.pkl")
```

Fitting 5 folds for each of 20 candidates, totalling 100 fits

Total execution time: 451.12 seconds

Best Params: {'ridge__alpha': 0.0006951927961775605}

Best Score: 0.011389699496520364

```
[16]: # Evaluate on the test set
      y_te_pred_log = model_ridge.predict(X_te)
```

```

ridge_model_name = "ridge regression"
ridge_mae_log, ridge_rmse_log, ridge_r2_log, ridge_mae_real, ridge_rmse_real,
↳ridge_r2_real = compute_statistics(ridge_model_name, y_te, y_te_pred_log)

```

```

ridge regression (log space):
  MAE : 0.0118
  RMSE: 0.0157
  R2 : 0.9999

```

```

ridge regression (real space):
  MAE : 1.6656
  RMSE: 5.2803
  R2 : 0.9999

```

6 5) Usage: predicting boardings for a stop

6.0.1 Observation :

We build an example of how to use the model to answer the following question:

“How many boardings will there be on 20/06/2025 at stop RIVE00 on line 12?”

6.0.2 Action :

We need to build the data row in the same format as the **X_tr** dataset in order to feed it into the model.

```
[17]: # X_tr.info()
```

```

[18]: from datetime import timedelta

# Colonnes attendues par ton modèle (ordre fixe)
REQ_COLS = [
    "ligne", "ligne_type_act", "horaire_type", "arret_code_long",
    "indice_semaine", "indice_jour_semaine", "annee", "mois",
    "log_descentes", "log_frequentation_totale", "log_delta_montees_descentes",
    "type_vehicule", "weather_code",
    "log_montees_lag_1", "log_desc_lag_1", "log_montees_lag_7", "log_desc_lag_7",
    "log_rolling_montees_7d", "log_rolling_desc_7d",
    "est_weekend", "est_ferie"
]

def make_X_pred(arret, ligne):
    # TPG_20250620 contient les données de tous les arrêts et ligne pour la
    ↳ journée du 20 juin.
    # il faut donc filtrer par arrêt et ligne

```

```

Next_TPG = TPG_20250620[(TPG_20250620["arret_code_long"] == arret) &
↳ (TPG_20250620["ligne"] == ligne)]

row = {
    "ligne": Next_TPG["ligne"].iloc[0],
    "ligne_type_act": Next_TPG["ligne_type_act"].iloc[0],
    "horaire_type": Next_TPG["horaire_type"].iloc[0],
    "arret_code_long": Next_TPG["arret_code_long"].iloc[0],

    "indice_semaine": Next_TPG["indice_semaine"].iloc[0],
    "indice_jour_semaine": Next_TPG["indice_jour_semaine"].iloc[0],
    "annee": Next_TPG["annee"].iloc[0],
    "mois": Next_TPG["mois"].iloc[0],

    "log_descentes": Next_TPG["log_descentes"].iloc[0],
    "log_frequentation_totale": Next_TPG["log_frequentation_totale"].
↳ iloc[0],
    "log_delta_montees_descentes": Next_TPG["log_delta_montees_descentes"].
↳ iloc[0],

    "type_vehicule": Next_TPG["type_vehicule"].iloc[0],
    "weather_code": Next_TPG["weather_code"].iloc[0],

    "log_montees_lag_1": Next_TPG["log_montees_lag_1"].iloc[0],
    "log_desc_lag_1": Next_TPG["log_desc_lag_1"].iloc[0],
    "log_montees_lag_7": Next_TPG["log_montees_lag_7"].iloc[0],
    "log_desc_lag_7": Next_TPG["log_desc_lag_7"].iloc[0],
    "log_rolling_montees_7d": Next_TPG["log_rolling_montees_7d"].iloc[0],
    "log_rolling_desc_7d": Next_TPG["log_rolling_desc_7d"].iloc[0],

    "est_weekend": Next_TPG["est_weekend"].iloc[0],
    "est_ferie": Next_TPG["est_ferie"].iloc[0]
}

X_pred = pd.DataFrame([[row.get(c, np.nan) for c in REQ_COLS]],
↳ columns=REQ_COLS)
return X_pred

```

```

[19]: # Génération de X_pred pour RIVE00 / ligne 12
arret_pred = "RIVE00"
ligne_pred = "12"
X_pred = make_X_pred(arret = arret_pred, ligne= ligne_pred)

display(X_pred.tail(7))
display(X_pred.shape)

```

```
# Prédiction
y_pred_log = model_ridge.predict(X_pred)[0]
y_pred = np.exp(y_pred_log)
print("y_pred_log:", y_pred_log)
print("exp(y_pred_log):", np.exp(y_pred_log))

print(f"Prediction (RIVE00, L12 ) : {y_pred:,.0f} montées")
```

```

ligne ligne_type_act horaire_type arret_code_long indice_semaine \
0 12 PRINCIPAL NORMAL RIVE00 25

indice_jour_semaine annee mois log_descentes log_frequentation_totale \
0 5 2025 6 8.104788 16.212851

... type_vehicule weather_code log_montees_lag_1 log_desc_lag_1 \
0 ... Tram 3 7.897572 8.046082

log_montees_lag_7 log_desc_lag_7 log_rolling_montees_7d \
0 7.993515 8.118806 7.820945

log_rolling_desc_7d est_weekend est_ferie
0 7.901882 1 0

[1 rows x 21 columns]

(1, 21)

y_pred_log: 8.117893653279854
exp(y_pred_log): 3353.9487207395387
Prediction (RIVE00, L12 ) : 3,354 montées
```

7 6) Verification

```
[20]: from math import fabs
Next_TPG = TPG_20250620[(TPG_20250620["arret_code_long"] == arret_pred) &
↳ (TPG_20250620["ligne"] == ligne_pred)]
y_real = Next_TPG["nb_de_montees"].iloc[0]
y_pred

abs_err = fabs(y_pred - y_real)
mape = abs_err / y_real * 100

print(f"Question: boardings at RIVE00 (L12) on 20/06/2025?")
print(f"Model prediction : {y_pred:,.2f}")
print(f"Actual value : {y_real:,.2f}")
print(f"Absolute error : {abs_err:,.2f} ({mape:,.2f} %)")
```

Question: boardings at RIVE00 (L12) on 20/06/2025?

Model prediction : 3,353.95
Actual value : 3,320.14
Absolute error : 33.81 (1.02 %)

La prediction nous donne 3354 montées contre 3320 pour le réel. Soit un écart de 1%

[21]: # 7) Conclusion

This case study illustrates how to use a regression model to predict the number of boardings at a specific stop, on a given date. In our example (RIVE00, line 12, June 20, 2025), the prediction is close to the actual value, highlighting the model's ability to provide relevant estimates.

Although the approach remains simplified, it clearly demonstrates the practical application of the model trained on TPG data.

[22]: #end